

Synthetic Population Catalyst

Table of contents

| | |
|--|-----------|
| Introduction | 5 |
| I Using SPC | 6 |
| 1 Getting started | 7 |
| 1.1 What SPC does | 7 |
| 1.2 What SPC outputs | 7 |
| 2 SPC Outputs | 8 |
| 2.1 Citing | 8 |
| 2.2 Versioning | 8 |
| 3 Outputs for England (Counties) | 10 |
| 3.1 Citing | 16 |
| 4 Outputs for Wales (ITL regions) | 17 |
| 4.1 Citing | 19 |
| 5 Outputs for Scotland (Police Divisions) | 20 |
| 5.1 Citing | 20 |
| 6 Using the SPC output file | 21 |
| 6.1 Javascript | 21 |
| 6.2 Python | 21 |
| 6.2.1 Converting to Pandas data-frames and CSV | 22 |
| 6.2.2 Converting .pb file to JSON format | 22 |
| 6.2.3 Converting to numpy arrays | 22 |
| 6.2.4 Visualizing venues | 23 |
| 7 Installation | 24 |
| 7.1 Dependencies | 24 |
| 7.2 Compiling SPC | 24 |
| 7.3 Troubleshooting downloading | 24 |
| 8 Creating new study areas | 25 |
| 8.1 Specifying the area | 25 |

| | | |
|-----------|---|-----------|
| 8.2 | Run SPC for the new area | 25 |
| 8.3 | (Optional) run SPC for lots of areas | 25 |
| 8.4 | Using the output | 26 |
| II | Understanding SPC | 27 |
| 9 | Introduction | 28 |
| 10 | Technical overview | 29 |
| 10.1 | Phase 1: Data preparation | 30 |
| 10.1.1 | SPENSER | 30 |
| 10.1.2 | Downloading and preparation of public data from various sources | 31 |
| 10.1.3 | Enriching SPENSER | 32 |
| 10.1.4 | Azure upload | 32 |
| 10.2 | Phase 2: Running SPC for a specific study area | 32 |
| 11 | Modelling methods | 33 |
| 11.1 | SPENSER and QUANT | 33 |
| 11.2 | BMI estimation | 33 |
| 11.3 | Income data | 34 |
| 11.3.1 | Methods | 34 |
| 11.3.2 | Comparison to reference values from ONS | 35 |
| 11.4 | Commuting flows | 38 |
| 11.4.1 | List of all workplaces in GB | 38 |
| 11.4.2 | Usage inside SPC | 39 |
| 12 | Data schema | 40 |
| 12.1 | Understanding the schema | 40 |
| 12.2 | Flows: modelling daily activites | 40 |
| 12.3 | Flow weights | 41 |
| 13 | Data sources | 42 |
| 13.1 | Utility data | 42 |
| | lookUp-GB.csv.gz | 42 |
| 13.2 | County level data | 43 |
| | pop_.csv.gz | 43 |
| 13.3 | National data | 45 |
| | businessRegistry.csv.gz | 45 |
| | GIS/ | 45 |
| | QUANT_RAMP_spc.tar.gz | 45 |
| | timeAtHomeIncreaseCTY.csv.gz | 46 |
| | diariesRef.csv.gz | 46 |

| | | |
|------------|---|-----------|
| III | Advanced | 47 |
| 14 | Developer guide | 48 |
| 14.1 | Updating the docs | 48 |
| 14.2 | Code hygiene | 48 |
| 14.3 | Some tips for working with Rust | 48 |
| 14.4 | Docker | 49 |
| 15 | Code walkthrough | 50 |
| 15.1 | Generally useful techniques | 50 |
| 15.1.1 | Split code into two stages | 50 |
| 15.1.2 | Explicit data schema | 50 |
| 15.1.3 | Type-safe IDs | 51 |
| 15.1.4 | Idempotent data preparation | 51 |
| 15.1.5 | Logging with structure | 52 |
| 15.1.6 | Determinism | 52 |
| 15.2 | Protocol buffers | 53 |
| 15.3 | An example of the power of static type checking | 53 |
| 16 | Performance | 55 |

Introduction



The Synthetic Population Catalyst (SPC) makes it easier for researchers to work with synthetic population data in Great Britain. It combines a variety of [data sources](#) and outputs a single file in [protocol buffer format](#), describing the population and its activities in a given study area. The data include socio-demographic, health, salary and daily activity data per person, and information about the venues where people conduct those activities.

SPC outputs can be used to catalyze other projects. Rather than join together many [raw data sources](#) yourself and deal with missing and messy data, you can leverage SPC's effort and well-documented schema.

You can download this site as [a PDF](#) and find all code [on Github](#).

This work was supported by Wave 1 of The UKRI Strategic Priorities Fund under the EP-SRC Grant EP/W006022/1, particularly the “Ecosystem of Digital Twins” and “Shocks and Resilience” themes within that grant & The Alan Turing Institute

Part I

Using SPC

1 Getting started

We suggest to start by exploring one of the pre-compiled areas we have made readily available:

1. Download [sample data](#) for an area in Great Britain
2. Unpack it and open it with the [web explorer](#)

Possible next steps:

3. Learn more on [how to use the data](#)
4. If you need a custom area, [build](#) and then [run](#) SPC

1.1 What SPC does

SPC generates spatially enriched synthetic population outputs for any area that is comprised of one or more Middle-Layer Output Areas (MSOAs) in England and Wales and/or one or more Intermediary Zones (IZ) in Scotland, including Local Authority Districts - LADs. The output file generated by SPC has a granularity of Output Area (150 ± 100 households). This file is structured to help other researchers or urban analysts to feed dynamic models, such as ABMs, for multiple purposes where an enriched synthetic population file is required. SPC includes a comprehensive set of variables that include sociodemographic characteristics, daily activities, and other extra data to help you model the complexity of British society.

1.2 What SPC outputs

You can see all of the per-person, household, and OA information SPC provides in the [schema](#) and [data sources](#). We use [protocol buffers](#) to efficiently encode the data and describe its shape.

2 SPC Outputs

You don't need to run SPC yourself. See [config/](#) for the list of MSOAs covered by each study area. If you want to run SPC for a different list of MSOAs, [see here](#).

One of the advantages of using SPC is that help researches to mimic the population characteristics and its iterations through multiples years (see for more details). So you can replicate what the population might look like across multiple periods of time. Initially check what country you would like to explore, then pick the year to get the outcome file. In case you want to explore it and see how does the data look like, and what attributes are included, load the output in our [SPC Explorer](#) and get inspired about the potential applications you could co-create using these outcomes.

- [England](#) (Available years: 2012, 2020, 2022, 2039)
- [Wales](#) (Available years: 2012, 2020, 2022, 2039)

We also included some special areas for your testing:

- [special/2012/northwest_transpennine.pb.gz](#)
- [special/2020/northwest_transpennine.pb.gz](#)
- [special/2022/northwest_transpennine.pb.gz](#)
- [special/2032/northwest_transpennine.pb.gz](#)
- [special/2039/northwest_transpennine.pb.gz](#)

2.1 Citing

If you use SPC code or data in your work, please cite using the [Zenodo DOI](#) (using the bottom-right tool to generate the citation).

2.2 Versioning

Over time, we may add more data to SPC or change the schema. Protocol buffers are designed to let combinations of new/old code and data files work together, but we don't intend to use this feature. We may make breaking changes, like deleting fields. We'll release a new version of the schema and output data every time and document it here. You should depend on a

specific version of the data output in your code, so new releases don't affect you until you decide to update.

- v1: released 25/04/2022, [schema](#)
- v1.1, released 27/05/2022, [schema](#)
 - added `pwkstat`, `salary_hourly`, `salary_yearly`, and `idp`
 - reorganized `Identifiers` and `Employment` attributes
 - non-breaking change added 02/08/2022: added `bmi_new` field
- v1.2, released 29/12/2022, [schema](#)
 - switched to `proto2` and made some fields optional
 - adjusted some numeric enum values to match ONS
- v2, released 09/03/2023, [schema](#)
 - new per-person and per-household fields
 - various changes to existing fields (adjusting enum number, removing the BMI enum, etc)
 - adding time-use diaries
 - expanding to Wales
 - adding multiple years of output

3 Outputs for England (Counties)

The counties of England are in this context the lieutenancy areas, often referred to as ceremonial counties. There are officially 48 of them, although we have chosen to include the City of London within Greater London in our release. Check the year you would like to explore and pick the corresponding file based on the region you are interested. Remember if you want to explore the data you can load the output in our [SPC explorer](#)

- 2012:
 - [bedfordshire.pb.gz](#)
 - [berkshire.pb.gz](#)
 - [bristol.pb.gz](#)
 - [buckinghamshire.pb.gz](#)
 - [cambridgeshire.pb.gz](#)
 - [cheshire.pb.gz](#)
 - [cornwall.pb.gz](#)
 - [cumbria.pb.gz](#)
 - [England/2012/derbyshire.pb.gz](#)
 - [devon.pb.gz](#)
 - [durham.pb.gz](#)
 - [east-sussex.pb.gz](#)
 - [east-yorkshire-with-hull.pb.gz](#)
 - [essex.pb.gz](#)
 - [gloucestershire.pb.gz](#)
 - [greater-london.pb.gz](#)
 - [greater-manchester.pb.gz](#)
 - [hampshire.pb.gz](#)
 - [herefordshire.pb.gz](#)
 - [hertfordshire.pb.gz](#)
 - [isle-of-wight.pb.gz](#)
 - [kent.pb.gz](#)
 - [lancashire.pb.gz](#)
 - [leicestershire.pb.gz](#)
 - [lincolnshire.pb.gz](#)
 - [merseyside.pb.gz](#)
 - [norfolk.pb.gz](#)

- northamptonshire.pb.gz
- northumberland.pb.gz
- north-yorkshire.pb.gz
- nottinghamshire.pb.gz
- oxfordshire.pb.gz
- rutland.pb.gz
- shropshire.pb.gz
- somerset.pb.gz
- south-yorkshire.pb.gz
- staffordshire.pb.gz
- suffolk.pb.gz
- surrey.pb.gz
- tyne-and-wear.pb.gz
- warwickshire.pb.gz
- west-midlands.pb.gz
- west-sussex.pb.gz
- west-yorkshire.pb.gz
- wiltshire.pb.gz
- worcestershire.pb.gz

- 2020:

- bedfordshire.pb.gz
- berkshire.pb.gz
- bristol.pb.gz
- buckinghamshire.pb.gz
- cambridgeshire.pb.gz
- cheshire.pb.gz
- cornwall.pb.gz
- cumbria.pb.gz
- derbyshire.pb.gz
- devon.pb.gz
- durham.pb.gz
- east-sussex.pb.gz
- east-yorkshire-with-hull.pb.gz
- essex.pb.gz
- gloucestershire.pb.gz
- greater-london.pb.gz
- greater-manchester.pb.gz
- hampshire.pb.gz
- herefordshire.pb.gz
- hertfordshire.pb.gz
- isle-of-wight.pb.gz

- kent.pb.gz
- lancashire.pb.gz
- leicestershire.pb.gz
- lincolnshire.pb.gz
- merseyside.pb.gz
- norfolk.pb.gz
- northamptonshire.pb.gz
- northumberland.pb.gz
- north-yorkshire.pb.gz
- nottinghamshire.pb.gz
- oxfordshire.pb.gz
- rutland.pb.gz
- shropshire.pb.gz
- somerset.pb.gz
- south-yorkshire.pb.gz
- staffordshire.pb.gz
- suffolk.pb.gz
- surrey.pb.gz
- tyne-and-wear.pb.gz
- warwickshire.pb.gz
- west-midlands.pb.gz
- west-sussex.pb.gz
- west-yorkshire.pb.gz
- wiltshire.pb.gz
- worcestershire.pb.gz

- 2022:

- bedfordshire.pb.gz
- berkshire.pb.gz
- bristol.pb.gz
- buckinghamshire.pb.gz
- cambridgeshire.pb.gz
- cheshire.pb.gz
- cornwall.pb.gz
- cumbria.pb.gz
- derbyshire.pb.gz
- devon.pb.gz
- durham.pb.gz
- east-sussex.pb.gz
- east-yorkshire-with-hull.pb.gz
- essex.pb.gz
- gloucestershire.pb.gz

- greater-london.pb.gz
- greater-manchester.pb.gz
- hampshire.pb.gz
- herefordshire.pb.gz
- hertfordshire.pb.gz
- isle-of-wight.pb.gz
- kent.pb.gz
- lancashire.pb.gz
- leicestershire.pb.gz
- lincolnshire.pb.gz
- merseyside.pb.gz
- norfolk.pb.gz
- northamptonshire.pb.gz
- northumberland.pb.gz
- north-yorkshire.pb.gz
- England/2022/nottinghamshire.pb.gz
- oxfordshire.pb.gz
- rutland.pb.gz
- shropshire.pb.gz
- somerset.pb.gz
- south-yorkshire.pb.gz
- staffordshire.pb.gz
- suffolk.pb.gz
- surrey.pb.gz
- tyne-and-wear.pb.gz
- warwickshire.pb.gz
- west-midlands.pb.gz
- west-sussex.pb.gz
- west-yorkshire.pb.gz
- wiltshire.pb.gz
- worcestershire.pb.gz

- 2032:

- bedfordshire.pb.gz
- berkshire.pb.gz
- bristol.pb.gz
- buckinghamshire.pb.gz
- cambridgeshire.pb.gz
- cheshire.pb.gz
- cornwall.pb.gz
- cumbria.pb.gz
- derbyshire.pb.gz

- devon.pb.gz
- durham.pb.gz
- east-sussex.pb.gz
- east-yorkshire-with-hull.pb.gz
- essex.pb.gz
- gloucestershire.pb.gz
- greater-london.pb.gz
- greater-manchester.pb.gz
- hampshire.pb.gz
- herefordshire.pb.gz
- hertfordshire.pb.gz
- isle-of-wight.pb.gz
- kent.pb.gz
- lancashire.pb.gz
- leicestershire.pb.gz
- lincolnshire.pb.gz
- merseyside.pb.gz
- norfolk.pb.gz
- northamptonshire.pb.gz
- northumberland.pb.gz
- north-yorkshire.pb.gz
- nottinghamshire.pb.gz
- oxfordshire.pb.gz
- rutland.pb.gz
- shropshire.pb.gz
- somerset.pb.gz
- south-yorkshire.pb.gz
- staffordshire.pb.gz
- suffolk.pb.gz
- surrey.pb.gz
- tyne-and-wear.pb.gz
- warwickshire.pb.gz
- west-midlands.pb.gz
- west-sussex.pb.gz
- west-yorkshire.pb.gz
- wiltshire.pb.gz
- worcestershire.pb.gz

- 2039:

- bedfordshire.pb.gz
- berkshire.pb.gz
- bristol.pb.gz

- buckinghamshire.pb.gz
- cambridgeshire.pb.gz
- cheshire.pb.gz
- cornwall.pb.gz
- cumbria.pb.gz
- derbyshire.pb.gz
- devon.pb.gz
- durham.pb.gz
- east-sussex.pb.gz
- east-yorkshire-with-hull.pb.gz
- essex.pb.gz
- gloucestershire.pb.gz
- greater-london.pb.gz
- greater-manchester.pb.gz
- hampshire.pb.gz
- herefordshire.pb.gz
- hertfordshire.pb.gz
- isle-of-wight.pb.gz
- kent.pb.gz
- lancashire.pb.gz
- leicestershire.pb.gz
- lincolnshire.pb.gz
- merseyside.pb.gz
- norfolk.pb.gz
- northamptonshire.pb.gz
- northumberland.pb.gz
- north-yorkshire.pb.gz
- nottinghamshire.pb.gz
- oxfordshire.pb.gz
- rutland.pb.gz
- shropshire.pb.gz
- somerset.pb.gz
- south-yorkshire.pb.gz
- staffordshire.pb.gz
- suffolk.pb.gz
- surrey.pb.gz
- tyne-and-wear.pb.gz
- warwickshire.pb.gz
- west-midlands.pb.gz
- west-sussex.pb.gz
- west-yorkshire.pb.gz
- wiltshire.pb.gz
- worcestershire.pb.gz

3.1 Citing

If you use SPC code or data in your work, please cite using the [Zenodo DOI](#) (using the bottom-right tool to generate the citation).

4 Outputs for Wales (ITL regions)

International Territorial Level (ITL) regions are a post-brexite renaming of the former Nomenclature of Territorial Units for Statistics (NUTS) regions. In wales, the level 3 represents a grouping of the 22 unitary districts into 12 regions. Check the year you would like to explore and pick the corresponding file based on the region you are interested. Remember if you want to explore the data you can load the output in our [SPC explorer](#)

- 2012:
 - [bridgend-and-neath-port-talbot.pb.gz](#)
 - [cardiff-and-vale-of-glamorgan.pb.gz](#)
 - [central-valleys.pb.gz](#)
 - [conwy-and-denbighshire.pb.gz](#)
 - [flintshire-and-wrexham.pb.gz](#)
 - [gwent-valleys.pb.gz](#)
 - [gwynedd.pb.gz](#)
 - [isle-of-anglesey.pb.gz](#)
 - [monmouthshire-and-newport.pb.gz](#)
 - [powys.pb.gz](#)
 - [south-west-wales.pb.gz](#)
 - [swansea.pb.gz](#)
- 2020:
 - [bridgend-and-neath-port-talbot.pb.gz](#)
 - [cardiff-and-vale-of-glamorgan.pb.gz](#)
 - [central-valleys.pb.gz](#)
 - [conwy-and-denbighshire.pb.gz](#)
 - [flintshire-and-wrexham.pb.gz](#)
 - [gwent-valleys.pb.gz](#)
 - [gwynedd.pb.gz](#)
 - [isle-of-anglesey.pb.gz](#)
 - [monmouthshire-and-newport.pb.gz](#)
 - [powys.pb.gz](#)
 - [south-west-wales.pb.gz](#)
 - [Wales/2020/swansea.pb.gz](#)
- 2022:

- [Wales/2022/bridgend-and-neath-port-talbot.pb.gz](#)
- [Wales/2022/cardiff-and-vale-of-glamorgan.pb.gz](#)
- [Wales/2022/central-valleys.pb.gz](#)
- [Wales/2022/conwy-and-denbighshire.pb.gz](#)
- [Wales/2022/flintshire-and-wrexham.pb.gz](#)
- [Wales/2022/gwent-valleys.pb.gz](#)
- [Wales/2022/gwynedd.pb.gz](#)
- [Wales/2022/isle-of-anglesey.pb.gz](#)
- [Wales/2022/monmouthshire-and-newport.pb.gz](#)
- [Wales/2022/powys.pb.gz](#)
- [Wales/2022/south-west-wales.pb.gz](#)
- [Wales/2022/swansea.pb.gz](#)

- 2032:

- [bridgend-and-neath-port-talbot.pb.gz](#)
- [cardiff-and-vale-of-glamorgan.pb.gz](#)
- [central-valleys.pb.gz](#)
- [conwy-and-denbighshire.pb.gz](#)
- [flintshire-and-wrexham.pb.gz](#)
- [gwent-valleys.pb.gz](#)
- [gwynedd.pb.gz](#)
- [isle-of-anglesey.pb.gz](#)
- [monmouthshire-and-newport.pb.gz](#)
- [powys.pb.gz](#)
- [south-west-wales.pb.gz](#)
- [swansea.pb.gz](#)

- 2039:

- [bridgend-and-neath-port-talbot.pb.gz](#)
- [cardiff-and-vale-of-glamorgan.pb.gz](#)
- [central-valleys.pb.gz](#)
- [conwy-and-denbighshire.pb.gz](#)
- [flintshire-and-wrexham.pb.gz](#)
- [gwent-valleys.pb.gz](#)
- [gwynedd.pb.gz](#)
- [isle-of-anglesey.pb.gz](#)
- [monmouthshire-and-newport.pb.gz](#)
- [powys.pb.gz](#)
- [south-west-wales.pb.gz](#)
- [swansea.pb.gz](#)

4.1 Citing

If you use SPC code or data in your work, please cite using the [Zenodo DOI](#) (using the bottom-right tool to generate the citation).

5 Outputs for Scotland (Police Divisions)

Police divisions are a convenient grouping of unitary districts. Check the year you would like to explore and pick the corresponding file based on the region you are interested. Remember if you want to explore the data you can load the output in our [SPC explorer](#)

- 2012:
 - TBC soon
- 2020:
 - TBC soon
- 2022:
 - TBC soon
- 2032:
 - TBC soon
- 2039:
 - TBC soon

5.1 Citing

If you use SPC code or data in your work, please cite using the [Zenodo DOI](#) (using the bottom-right tool to generate the citation).

6 Using the SPC output file

Once you [download](#) or [generate](#) an SPC output file for your study area, how do you use it? Each study area consists of one `.pb` or [protocol buffer file](#). This file efficiently encodes data following this [schema](#). [Read more](#) about what data is contained in the output.

You can read the “protobuf” (shorthand for a protocol buffer file) in any [supported language](#), and then extract and transform just the parts of the data you want for your model.

We have examples for Python below, but feel free to request other languages.

6.1 Javascript

We have a [web app](#) using Svelte to interactively explore SPC data. Its [source code](#) is great reference for how to use the proto output.

6.2 Python

To work with SPC protobufs in Python, you need two dependencies setup:

- The [protobuf](#) library
 - You can install system-wide with `pip install protobuf`
 - Or add as a dependency to a conda, poetry, etc environment
- The generated Python library, [synthpop_pb2.py](#)
 - You can download a copy of this file into your codebase, then `import synthpop_pb2`
 - You can also generate the file yourself, following the [docs](#): `protoc --python_out=python/synthpop.proto`

6.2.1 Converting to Pandas data-frames and CSV

The [schema](#) expresses relationships between people, households, and venues that can't all be captured by a simple 2D table. Nevertheless, you can extract per-person information and express as a dataframe or CSV file. See [this example Python script](#) for inspiration. You can try it out:

```
# Download a file
wget https://ramp0storage.blob.core.windows.net/spc-output/v1/rutland.pb.gz
# Uncompress
gunzip rutland.pb.gz
# Convert the .pb to JSON
python3 python/protobuf_to_csv.py --input_path data/output/rutland.pb
# View the output
less people.csv
```

6.2.2 Converting .pb file to JSON format

To interactively explore the data, viewing JSON is much easier. It shows the same structure as the protobuf, but in a human-readable text format. The example below uses a [small Python script](#):

```
# Download a file
wget https://ramp0storage.blob.core.windows.net/spc-output/v1/rutland.pb.gz
# Uncompress
gunzip rutland.pb.gz
# Convert the .pb to JSON
python3 python/protobuf_to_json.py data/output/rutland.pb > rutland.json
# View the output
less rutland.json
```

6.2.3 Converting to numpy arrays

The [ASPICS](#) project simulates the spread of COVID through a population. The code uses numpy, and [this script](#) converts the protobuf to a bunch of different numpy arrays.

Note the ASPICS code doesn't keep using the generated Python protobuf classes for the rest of the pipeline. Data frames and numpy arrays may be more familiar and appropriate. The protobuf is a format optimized for reading and writing; you don't need to use it throughout all of your model code.

6.2.4 Visualizing venues

Use [this script](#) to read a protobuf file, then draws a dot for every venue, color-coded by activity.



7 Installation

You only need to compile SPC to run for a custom set of MSOAs. Just [download existing output](#) if your study area matches what we provide.

7.1 Dependencies

- **Rust:** The latest stable version of Rust: <https://www.rust-lang.org/tools/install>

7.2 Compiling SPC

```
git clone https://github.com/alan-turing-institute/uatk-spc/  
cd uatk-spc  
# The next command will take a few minutes the first time you do it, to build external dep  
cargo build --release
```

7.3 Troubleshooting downloading

If you get an error `No such file or directory (os error 2)` it might be because a previous attempt to run SPC failed, and some necessary files were not fully downloaded. In these cases you could try deleting the `data/raw_data` directory and then running SPC again. It should automatically try to download the big files again.

If you have trouble downloading any of the large files, you can download them manually. The logs will contain a line such as `Downloading https://ramp0storage.blob.core.windows.net/nationaldata/` to `data/raw_data/nationaldata/QUANT_RAMP_spc.tar.gz`. This tells you the URL to retrieve, and where to put the output file. Note that SPC won't attempt to download files if they already exist, so if you wind up with a partially downloaded file, you have to manually remove it.

8 Creating new study areas

If the area you want to model isn't [already generated](#), then you can follow this guide to run SPC on a custom area. You must first [compile SPC](#).

8.1 Specifying the area

SPC takes a newline-separated list of MSOAs in the `config/` directory as input, like [this](#). You can generate this list from a LAD (local authority district). From the main SPC directory, run `python scripts/select_msoas.py`. Refer to `data/raw_data/referencedata/lookUp.csv` (only available after running SPC once) for all geographies available.

This script will create a new file, `config/your_region.txt`.

8.2 Run SPC for the new area

From the main directory, just run:

```
cargo run --release -- config/your_region.txt
```

This will download some large files the first time. You'll wind up with `data/output/your_region.pb` as output, as well as lots of intermediate files in `data/raw_data/`. The next time you run this command (even on a different study area), it should go much faster.

8.3 (Optional) run SPC for lots of areas

If you want to run the program over lots of areas at once and are using Mac/Linux, you can use a `for` loop in a terminal to repeatedly run SPC over all files in the `config` directory. For example, this will run SPC on all `.txt` files in the `config` directory:

```
for file in config/*.csv; do cargo run --release -- config/$file; done
```

8.4 Using the output

After you generate the files, see [here](#) for how to use them in your project.

If you use SPC code or data in your work, please cite using the [Zenodo DOI](#) (using the bottom-right tool to generate the citation).

Part II

Understanding SPC

9 Introduction

Blabla

10 Technical overview

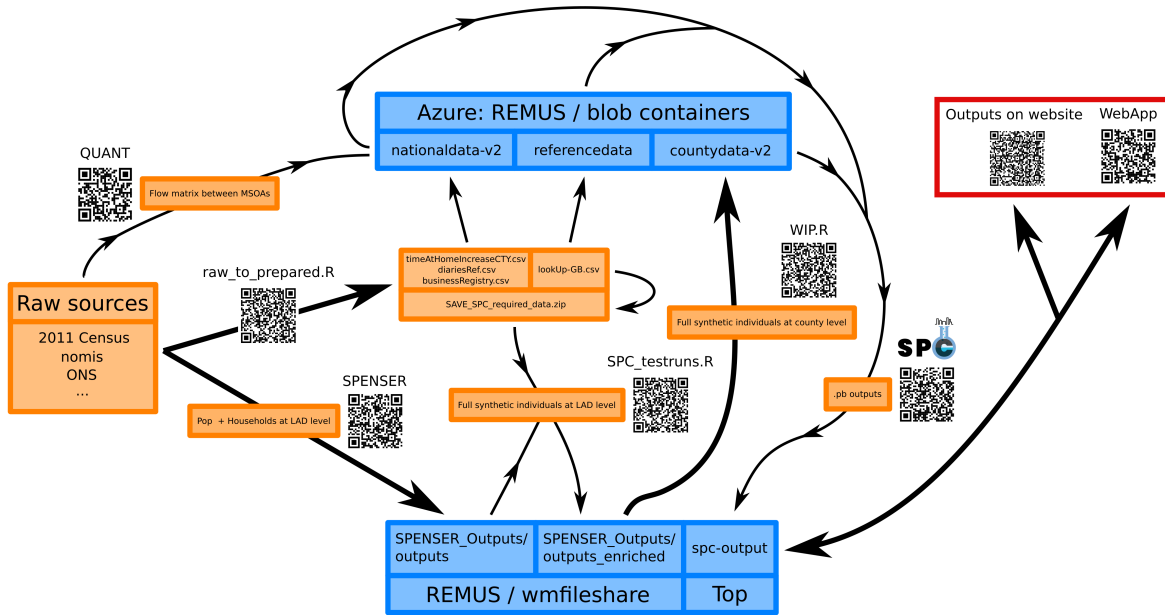
SPC is divided into two phases. The data preparation phase relies on scripts that only need to be run once. It is meant to output a postprocessed version of all the raw data sources that allows the model to run smoothly on custom areas. The second phase involves the user choosing a custom area and launching a simulation. It pulls the relevant datasets among the data prepared by the first phase, calculates the different daily activities and formats the results into a single protobuf file.

We provide in this document a step by step description of running the entire SPC pipeline. For

The full SPC pipeline comprises the following steps.

- **Phase 1** - Data preparation steps:
 1. The [SPENSER](#) model creates a synthetic population with basic demographic information for all of GB.
 2. A [script](#) downloads and prepares data from various public sources that will be used throughout the model.
 3. The outputs of SPENSER are enriched using some of the outputs from step 2.
 4. The resulting outputs are uploaded as `.csv` files to a dedicated Azure repository.
- **Phase 2** - Steps after the user has selected a study area:
 1. All the data relevant to the study area are pulled from Azure.
 2. Individuals are assigned a single education destination (if under 16) and several potential retail destinations, according to a local version of the [QUANT](#) model.
 3. Individuals are assigned a single workplace destination (if above 16), according to the method described [here](#).
 4. The population, its activities and an optional lockdown modelling are gathered into a single `.pb` file that can be visualised by the [SPC explorer](#).

We now explain how to run each step. The theoretical concepts supporting the modelling are presented [here](#).



10.1 Phase 1: Data preparation

10.1.1 SPENSER

The original [SPENSER](#) model is made up of 5 different GitHub repositories, operating specific parts of the simulation of a synthetic population (gathering the data from ONS, creating individuals, creating households, assigning individuals to households and projecting the population to future years). We use this [modified version](#) with instructions for running the full pipeline on a [single machine](#).

The SPENSER microsimulation is split into three steps:

1. **Household synthesis**: households are synthesised for a base year (2011) from [census data](#) at OA resolution. These households are then sequentially synthesised for subsequent years using [household forecasts](#).
2. **Population synthesis**: people are sequentially synthesised using marginal census data on gender, age and ethnicity at MSOA resolution for 2011, with population projections used to derive the marginals beyond the reference census year.
3. **Assignment**: for a given year, the synthesised population (from step 2) can be assigned to a synthesised household (from step 1), while a “[household representative person](#)” from the synthesised population (step 2) is assigned to each synthesised household (from step 1).

The result of SPENSER is two separate datasets and a merging key: one dataset for individuals, accurate at MSOA level and containing the `sex`, `age` and `ethnicity` fields; and one for households, accurate at OA level and containing the `OA11CD`, `HOUSE_nssec8`, `House_type`, `HOUSE_typeCommunal`, `HOUSE_NRooms`, `HOUSE_centralHeat`, `HOUSE_tenure` and `HOUSE_NCars` fields respectively.

10.1.2 Downloading and preparation of public data from various sources

Instructions to run this step from the source code can be found under [Step 1: Curate public data from diverse sources](#). The result is a set of data files, some of which will be merged with the outputs from SPENSER during the next step, containing:

- NSSEC8 distributions among the population of England and Wales by age group and sex at MSOA level (`NSSEC8_EW_F_16to24_CLEAN.csv`, etc.) and among the total population of Scotland by age group, sex and ethnicity (`NSSECS_CLEAN.csv`)
- A combined extract from the three latest GB Health Surveys (`HSComplete.csv`)
- An extract from the UK Time Use Survey 2015 (`indivTUS.csv`)
- A file containing a set of coefficients to estimate the average BMI of individuals in England depending on their age, sex and ethnicity (`BMIdMean.csv`) and a file containing coefficients to obtain the equivalent average BMI in Scotland and Wales (`BMIdiff.csv`)
- Coefficients to estimate the hourly salary of an employee in England depending on their home region, sex, part-time/full-time status, age and SOC category (`coefFFT.csv`, etc. & `ageRescaleFFT.csv`, etc.).
- Coefficients to estimate the numbers of hours worked corresponding to the criteria mentioned above (`meanHoursFFT.csv`, etc. and `distribHours.csv`)
- Centroid coordinates of Output areas in GB (`OACentroids.csv`)

In addition, four files to be used by the second phase of the model are outputted:

- `diariesRef.csv` contains diaries of typical days extracted from the UK Time Use Survey
- `businessRegistry.csv` contains a list of all individual workplaces in GB
- `timeAtHomeIncreaseCTY.csv` contains a reduction in time spent away from home during the pandemic according to Google Mobility reports
- `lookUp-GB.csv` is a comprehensive lookup table between GB geographies, including name variants used by Google and OSM and local file names for storage within Azure

To understand the methods supporting the creation of these files, we refer the reader to the corresponding section inside the [modelling methods](#) section.

10.1.3 Enriching SPENSER

Instructions to run this step can be found under [Step 2: Add to SPENSER](#). Line numbers quoted in the following refer to this [script](#).

Once merged into one dataset according to the matching key (l. 13-49), the SPENSER data is enriched with the outputs of the previous step. An individual among those sharing the same 5-year age group (extra details for under 18) and sex is drawn from the participants of the Health Survey (l. 56-72). This adds the `id_HS`, `HEALTH_diabetes`, `HEALTH_bloodpressure`, `HEALTH_cvd`, `HEALTH_NMedecines`, `HEALTH_selfAssessed` and `HEALTH_lifeSat` [fields](#). This join is not spatially differentiated and other matching criteria (ethnicity and `nssec8`) were retained due to a lack of representativity inside the survey. The BMI field is then added l. 74-89, according to [this method](#).

Each individual that is not a head of household is assigned an `nssec8` category (l. 96-108). This is done according to `nssec8` category distributions among the general population by sex and age groups according to ONS data ([DC6114EW](#) and [DC6206SC](#) datasets). An individual among those sharing the same 5-year age group, sex and `nssec8` category is drawn from the participants of the UK Time Use Survey (l. 111-125). This adds the `id_TUS_hh`, `id_TUS_p`, `pwkstat`, `soc2010`, `sic1d2007`, `sic2d2007`, `netPayWeekly` and `workedHoursWeekly` [fields](#). Note that the `netPayWeekly` and `workedHoursWeekly` fields have a low response rate among participants of the survey. For that reason, we have added a [much more detailed modelling of income](#), that includes spatial differences at region level (l. 130-140).

Coordinates of the OA where the household's home is located are finally added l. 152-156.

10.1.4 Azure upload

Following enrichment, a final step involves [grouping LADs into counties](#) and [uploading to an Azure container](#) for use as input for Phase 2 below.

10.2 Phase 2: Running SPC for a specific study area

This part is corresponding to the scripts written in Rust. Instructions can be found [here](#).

11 Modelling methods

We present here the theoretical principles behind the modelling done in SPC and point to the parts of the code where they are implemented.

11.1 SPENSER and QUANT

The generation of the population data by [SPENSER](#) and the modelling of trips to schools and retail by [QUANT](#) are detailed in

Lomax N et al. An Open-Source Model for Projecting Small Area Demographic and Land-Use Change. *Geographical analysis*, 54(3), 599-622 (2022). ([DOI](#))

and

Spooner F et al. A dynamic microsimulation model for epidemics. *Soc Sci Med.*, 291:114461 (2021). ([DOI](#))

11.2 BMI estimation

Body Mass Index (BMI) is calculated for each individual from the [Health Survey for England 2019](#) (access needs to be requested to the UK Data Service). This calculation is independent from the matching with the Health Survey that happens during the data preparation step, therefore the BMI values will not match the ones that could be obtained from the Health Survey identifiers included in the output. As the BMI variable is not necessarily independent from the other health variables (diabetes etc.), the new variable should only be used for studies where all other variables are considered equal. The new variable is continuous (a float) instead of categorical.

According to the HSE 2019, the distribution of BMI values should follow figure 1. The socioeconomic category variable was discarded for the modelling as it is not independent from the other variables, and “mixed” and “other” ethnicity categories have been merged due to small sample sizes.

Figure 1. BMI per age. Columns represent ethnicity (White, Black, Asian, Other), and the rows sex (female, male).

The distribution for each age group is a gamma distribution. See figure 2.

Figure 2. Distribution of BMI values for white females aged 30-34.

Due to small sample sizes, the BMI is calculated for each individual depending on their age according to a gamma distribution whose mean is the mean for the corresponding age, sex and ethnicity (thick line in figure 1), but whose variance is only determined by the total variance by sex and ethnicity. The resulting BMI values were validated for Bedfordshire, and correlations of 0.93 and 0.97 were found between the mean and variance of the modelled data compared to those for the reference HSE 2019 data. See figure 3. The distribution per age, as in figure 1, were also validated.

Figure 3. Modelled mean and variance compared to the reference mean and variance from the HSE 2019 data for each of the eight categories of figure 1.

The R code for this modelling are l. 239-406 of this [script](#), and the validation is included in the legacy version, [here](#).

11.3 Income data

This modelling is based on the 2020 revised edition of the [Earnings and hours worked, region by occupation by four-digit SOC: ASHE Table 15](#) database from ONS. Some percentiles for employees' gross hourly salaries are provided for each full-time and part-time job according to their four-digit SOC classification per region, and separated by sex. It is supported by this [script](#).

11.3.1 Methods

The data are far from complete (only about 15% of all possible values), especially for the highest deciles. We found that the missing values amongst the partially filled SOCs could be estimated by interpolating an order 3 polynomial to the existing values. We found that the order 3 polynomials were a good fit for most categories (93.11%). SOCs with too many missing values are given the value for the category that is immediately higher in the SOC hierarchy. For some jobs, the highest percentiles seem capped, making the polynomial fitting fail. In that case, we have replaced the unknown values with the highest known value (as there is no clear and systematic fitting for these special cases). In addition, the highest decile is never detailed in the data, which means that the highest salaries are underestimated in the model (and exceptionally high salaries are not present). The result of this phase is four tables {male full-time, male part-time, female full-time, female part-time} containing the coefficients of the fitted order 3 polynomials, with an optional cap when relevant. This step is done in section 1 of the script.

A percentile is chosen randomly (uniformly) for each individual in England, and the salary is then deduced according to their full-time/part-time status, region, sex and SOC category. [Age data](#) from ONS are then integrated. Part of the differences that can be observed between different age groups is already taken into account through the SOC variable, since it is expected to evolve throughout an individual's career. To avoid counting this dependency twice, we compute the residual between the results of the initial modelling that does not take into account age and the expected results by age group according to the data. We then deduce a function that modifies a posteriori the estimated salary of an individual depending on their age, so that the salaries sum correctly by age groups. This step is done in sections 2 to 4 of the script.

To get the number of hours worked per week, we also use the ASHE Table 15. Since only minimal differences are observed between SOC categories, we simply complete missing values by approximating them by values of the category that is immediately higher within the SOC hierarchy. This step is done in section 5 of the script.

When added to the SPC population data, a basic hourly salary column is added, as well as a corresponding annual salary deduced from the number of worked hours. In addition, we repeat this process for all individuals that are categorised as 'Self-employed' or 'Employee unspecified' by the Time Use Survey matching,, as if they were full time employees. These values are recorded in the columns `IncomeHAsIF` and `IncomeYAsIf`.

11.3.2 Comparison to reference values from ONS

We compare the results of the modelling to the raw datasets from ONS.

- Mod for modelled
- M for male
- F for female
- H for hourly gross salary
- Y for annual gross salary
- FT for full-Time
- PT for part-Time
- Only individuals recorded as employees (i.e. not self-employed) are taken into account in this section.

Number of employees per sex and full-time/part-time classification

The numbers given by ONS vary from dataset to dataset and are reported by ONS as indicative only. For the modelled values, we give the total number of individuals with a non-zero salary in each category.

| Variable | All | FT | PT | M | M FT | M PT | F | F FT | F PT |
|----------|--------|--------|------|--------|-------|--------|--------|----------|----------|
| ONS tot | 22-26k | 16-19k | 6-8k | 11-13k | 9-11k | 1.5-2k | 11-13k | 6.5-7.5k | 4.5-5.5k |
| Mod tot | 23.1k | 18.5k | 4.6k | 11.8k | 11k | 0.8k | 11.3k | 7.5k | 3.8k |
| H | | | | | | | | | |
| Mod tot | 17.6k | 14.8k | 2.8k | 9.4k | 8.9k | 0.5k | 8.2k | 5.9k | 2.3k |
| Y | | | | | | | | | |

A significant number of individuals listed as working either full or part time have 0 effective worked hours per day according to the Time Use Survey matching. In those cases, an hourly salary is modelled depending on their SOC, region and sex, as for any other employee, but the annual salary will be displayed as 0. It is possible to estimate the likely true number of hours worked from the same ONS dataset (Table 15.9a: Paid hours worked - Total 2020), also depending on their sex, soc and region. This calculation has been added to the “As If” column.

Hourly gross salary per sex and full-time/part-time classification

| Variable | All | FT | PT | M | M FT | M PT | F | F FT | F PT |
|------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| ONS mean | 17.63 | 18.32 | 13.93 | 18.81 | 19.12 | 14.69 | 16.19 | 17.08 | 13.68 |
| ONS median | 13.71 | 15.15 | 10.38 | 14.84 | 15.58 | 10.12 | 12.58 | 14.42 | 10.47 |
| Mod mean | 16.45 | 17.19 | 13.45 | 17.50 | 17.84 | 12.75 | 15.35 | 16.23 | 13.60 |
| Mod median | 13.55 | 14.46 | 10.23 | 14.27 | 14.72 | 9.16 | 12.79 | 14.12 | 10.51 |

The median values are quite close to the ONS values, but the mean values are always lower. This is expected, see the description of the modelling above.

Annual gross salary per sex and full-time/part-time classification

Only values > 0 are retained for these calculations.

| Variable | All | FT | PT | M | M FT | M PT | F | F FT | F PT |
|------------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| ONS mean | 31,646 | 38,552 | 13,819 | 38,421 | 42,072 | 14,796 | 24,871 | 33,253 | 13,512 |
| ONS median | 25,886 | 31,487 | 11,240 | 31,393 | 33,915 | 10,883 | 20,614 | 28,002 | 4,743 |
| Mod mean | 34,317 | 36,595 | 22,257 | 37,574 | 38,496 | 20,698 | 30,594 | 33,729 | 22,585 |
| Mod median | 28,713 | 30,942 | 17,928 | 31,404 | 32,382 | 17,382 | 25,875 | 29,028 | 18,137 |

The average salary for part-time employees is correct when values equal to 0 are taken into account. This suggests that the total number of hours worked for part-time employees is

correct, but the way they are distributed among individuals is not. It could be due to the TUS taking a snapshot of the situation during a particular week, rather than averaging their data over the year. It appears that the TUS matching also overestimates the average number of hours worked for female employees.

Regional differences (hourly gross salary)

| Region | East East lands | East Mid- lands | London | North East | North West | South East | South West | West Mid- lands | Yorkshire and The Humber |
|------------|-----------------------|-----------------------|--------|---------------|---------------|---------------|---------------|-----------------------|-----------------------------|
| ONS mean | 16.74 | 15.87 | 23.78 | 15.69 | 16.36 | 17.88 | 16.36 | 16.34 | 15.76 |
| ONS median | 13.28 | 12.65 | 18.30 | 12.40 | 12.90 | 14.33 | 12.74 | 12.92 | 12.46 |
| Mod mean | 16.67 | 15.29 | 19.39 | 15.05 | 15.22 | 17.34 | 15.92 | 15.47 | 14.41 |
| Mod median | 13.69 | 12.79 | 16.25 | 12.42 | 12.44 | 14.84 | 13.35 | 12.64 | 12.44 |

The pearson correlations for mean and median between the modelled and raw values are 0.92 and 0.93.

Hourly gross salary per one-digit SOC

| 1d SOC | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |
|------------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| ONS mean | 26.77 | 23.38 | 18.29 | 13.42 | 13.35 | 10.87 | 10.94 | 12.23 | 10.77 |
| ONS median | 20.96 | 21.34 | 15.66 | 11.54 | 12.04 | 10.08 | 9.52 | 10.93 | 9.22 |
| Mod mean | 21.52 | 22.14 | 16.00 | 12.76 | 12.55 | 10.49 | 10.50 | 12.05 | 9.87 |
| Mod median | 17.22 | 20.66 | 14.12 | 11.46 | 11.34 | 9.71 | 9.59 | 10.82 | 9.12 |

1. Managers, directors and senior officials
2. Professional occupations
3. Associate professional and technical occupations
4. Administrative and secretarial occupations
5. Skilled trades occupations
6. Caring, leisure and other service occupations
7. Sales and customer service occupations
8. Process, plant and machine operatives
9. Elementary occupations.

The Pearson correlations for mean and median between the modelled and raw values are 0.98 and 0.98.

Hourly gross salary per age

The reference for this table is: [Table 6.5a Hourly pay - Gross 2020](#)

Table before weighting by age:

| Age | 16-17 | 18-21 | 22-29 | 30-39 | 40-49 | 50-59 | 60+ |
|------------|-------|-------|-------|-------|-------|-------|-------|
| ONS mean | 7.21 | 9.59 | 14.09 | 18.13 | 20.04 | 19.12 | 16.32 |
| ONS median | 6.36 | 9.00 | 12.26 | 15.08 | 15.89 | 14.39 | 12.17 |
| Mod mean | 12.77 | 14.96 | 16.33 | 16.93 | 16.83 | 16.66 | 16.29 |
| Mod median | 10.93 | 12.71 | 13.88 | 14.02 | 13.96 | 13.85 | 13.65 |

The Pearson correlations for mean and median between the modelled and raw values are 0.92 and 0.92.

Table after weighting by age:

| Age | 16-17 | 18-21 | 22-29 | 30-39 | 40-49 | 50-59 | 60+ |
|------------|-------|-------|-------|-------|-------|-------|-------|
| ONS mean | 7.21 | 9.59 | 14.09 | 18.13 | 20.04 | 19.12 | 16.32 |
| ONS median | 6.36 | 9.00 | 12.26 | 15.08 | 15.89 | 14.39 | 12.17 |
| Mod mean | 9.05 | 11.15 | 14.87 | 17.35 | 17.96 | 17.47 | 15.41 |
| Mod median | 8.20 | 9.51 | 12.86 | 14.41 | 14.78 | 14.43 | 12.56 |

The Pearson correlations for mean and median between the modelled and raw values are 0.99 and 0.99.

11.4 Commuting flows

11.4.1 List of all workplaces in GB

In order to distribute each individual of the population to a unique physical workplace, we first created a population of all individual workplaces in England, based on a combination of the Nomis UK Business Counts 2020 dataset and the Nomis Business register and Employment Survey 2015 (see [Data sources](#)). The first dataset gives the number of individual workplace counts per industry, using the SIC 2007 industry classification, with imprecise size (i.e. number of employees) bands at MSOA level. The second dataset gives the total number of jobs available at LSOA level per SIC 2007 industry category. We found that the distribution of workplace sizes follows closely a simple $1/x$ distribution, allowing us to draw for each workplace a size

within their band, with sum constraints given by the total number of jobs available, according to the second dataset. The R codes to create the list of all workplaces can be found [here](#).

11.4.2 Usage inside SPC

The workplace ‘population’ and individual population are levelled for each SIC 2007 category by removing the exceeding part of whichever dataset lists more items. This takes into account that people and business companies are likely to over-report their working availability (e.g. part time and seasonal contracts are not counted differently than full time contracts, job seekers or people on maternity leave might report the SIC of their last job). This process can be controlled by a threshold in the parameter file that defines the maximal total proportion of workers or jobs that can be removed. If the two datasets cannot be levelled accordingly, the categories are dropped and the datasets are levelled globally. Tests in the West Yorkshire area have shown that when the level 1 SIC, containing 21 unique categories, is used, 90% of the volume of commuting flows were recovered compared to the Nomis commuting OD matrices at MSOA level.

The employees for each workplace are drawn according to the ‘universal law of visitation’, see

Schläpfer M et al. The universal visitation law of human mobility. Nature 593, 522–527 (2021). ([DOI](#))

This framework predicts that visitors to any destination follow a simple

$$(r,f) = K / (rf)^2$$

distribution, where (r,f) is the density of visitors coming from a distance r with frequency f and K is a balancing constant depending on the specific area. In the context of commuting, it can be assumed that $f = 1$. Additionally, we only need to weigh potential employees against each other, which removes the necessity to compute explicitly K . In the West Yorkshire test, we found a Pearson coefficient of 0.7 between the predicted flows when aggregated at MSOA level and the OD matrix at MSOA level available from Nomis.

12 Data schema

12.1 Understanding the schema

Here are some helpful tips for understanding the [schema](#).

Each `.pb` file contains exactly one `Population` message. In contrast to datasets consisting of multiple `.csv` files, just a single file contains everything. Some of the fields in `Population` are lists (of people and households) or maps (of venues keyed by activity, or of MSOAs). Unlike a flat `.csv` table, there may be more lists embedded later. Each `Household` has a list of `members`, for example.

The different objects refer to each other, forming a graph structure. The protobuf uses `uint64` IDs to index into other lists. For example, if some household has `members = [3, 10]`, then those two people can be found at `population.people[3]` and `population.people[10]`. Each of them will have the same `household` ID, pointing back to something in the `population.households` list.

12.2 Flows: modelling daily activities

SPC models daily travel behavior of people as “flows.” Flows are broken down by by an [activity](#) – shopping/retail, attending primary or secondary school, working, or staying at home. For each activity type, a person has a list of venues where they may do that activity, weighted by a probability of going to that particular venue.

Note that `flows_per_activity` is stored in `InfoPerMSOA`, not `Person`. The flows for retail and school are only known at the MSOA level, not individually. So given a particular `Person` object, you first look up their household’s MSOA – `msoa = population.households[person.household].msoa` and then look up flows for that MSOA – `population.info_per_msoa[msoa].flows_per_activity`.

Each person has exactly 1 flow for home – it’s just `person.household` with probability 1. A person has 0 or 1 flows to work, based on the value of `person.workplace`.

This doesn’t mean that all people in the same MSOA share the same travel behavior. Each person has their own `activity_durations` field, based on time-use survey data. Even if two

people share the same set of places where they may go shopping, one person may spend much more time on that activity than another.

See the [ASPICS conversion script](#) for all of this in action – it has a function to collapse a person’s flows down into a single weighted list.

Note that per MSOA, very few venues are represented as destinations – 10 for retail and 5 for school. Only the most likely venues from QUANT are used.

12.3 Flow weights

How do you interpret the probabilities/weights for flows? If your model needs people to visit specific places each day, you could randomly sample a venue from the flows, weighting them appropriately. For retail, you may want to repeat this sampling every day of the simulation, so they visit different venues. For primary and secondary school, it may be more appropriate to sample once and store that for the simulation – a student probably doesn’t switch schools daily.

Alternatively, you can follow what ASPICS does. Every day, each person logically visits all possible venues, but their interaction there (possibly receiving or transmitting COVID) is weighted by the probability of each venue.

13 Data sources

The original data are provided at different scales, which define their level of accuracy. For simplicity, the outputs of SPC are geolocated at Output Area (OA) level, although this scale may not be relevant to all indicators. The 2011 OAs are a geographical unit created for census collection and are designed to be relatively homogeneous, with an average size between 120 and 129 households.

The data from Open Street Map (OSM) is downloaded directly from <https://www.openstreetmap.org>. Everything else is hosted through local copies inside one Azure repository that interacts automatically with the model. We describe below the content of this repository and indicate the raw source used for each indicator. It is divided into utilities, county level data and national data. To recreate the content of this repository from raw sources, please refer to [this part of the code](#).

13.1 Utility data

[lookUp-GB.csv.gz](#)

The look-up table links different geographies of Great Britain together. It is used internally by the model, but can also help the user define their own study area. The following are standard denominations, compatible with ONS fields of the same name. They are based on ONS [lookups](#). See ONS documentation for more details.

- OA11CD: Output area codes for the 2011 census (120 to 129 households)
- LSOA11CD & LSOA11NM: Lower-layer Super Output Areas (about 2000 individuals), replaced by Intermediary Zones for Scotland
- MSOA11CD,MSOA11NM: Middle-layer Super Output Areas (about 8000 individuals), replaced by Data Zones for Scotland
- LAD20CD, LAD20NM: Local Authority Districts (314 for England, 22 for Wales and 32 for Scotland)
- ITL321CD, ITL321NM, ITL221CD, ITL221NM, ITL121CD & ITL121NM: International Territorial Level, replacing pre-Brexit NUTS European divisions.
- RGN20CD & RGN20NM: Regions of England (NA for other Wales and Scotland)
- Country: England, Wales or Scotland

In addition,

- “AzureRef”: Name of the geographical unit for the County level data folder inside Azure (Lieutenancy Areas – a.k.a. Ceremonial Counties – for England, Scottish Police Divisions and ITL321NM for Wales) For Wales: ITL321NM
- “GoogleMob” & “OSM” are alternate spellings used by Google and OSM for their data releases.

13.2 County level data

Files in this section are grouped by country (England, Wales and Scotland), then date (2012, 2020, 2022, 2032, 2039). The format of a path to an individual file is:

```
https://ramp0storage.blob.core.windows.net/countydata-v2/[country]/[date]/pop_[area_name].csv
```

As of March 2023, England contains 5 series of 46 files (Dorset is missing), Wales 5 series of 12 files and Scotland is missing.

pop_.csv.gz

The data is mainly based on the [2011 UK census](#), the UK [Time Use Survey 2014-15](#) and the health surveys of GB ([England](#), [Wales](#), [Scotland](#)). The SPENSER (Synthetic Population Estimation and Scenario Projection) microsimulation model ([ref](#)) distributes individuals from the census with MSOA scale constraints into synthetic households with OA constraints. It is able to project this synthetic population in the future according to estimates from the Office for National Statistics (ONS). These data were enriched with some of the content of the other datasets mentioned (the rest of which can be added *a posteriori* from the identifiers provided). The data have also been complemented with a modelling of BMI and salaries. The methods used to join the different datasets are explained in the [methods](#).

The fields currently contained are detailed [here](#). They are:

- **pid**: Unique person identifier at GB level within SPC
- **hid**: Unique household identifier at GB level within SPC
- **OA11CD**: Output Area code of the individual’s home (ONS, 2011 boundaries)
- **sex**: Sex assigned at birth (DC1117EW, census 2011)
- **age**: Age in years (DC1117EW, census 2011)
- **ethnicity**: Based on self-report (aggregated from DC2101EW, census 2011)
- **nssec8**: National Statistics Socio-economic classification (see methods)
- **HOUSE_nssec8**: National Statistics Socio-economic classification of the reference person of the household (LC4605, census 2011)

- **House_type**: Type of accommodation (based on LC4402EW, census 2011)
- **HOUSE_typeCommunal**: Type of communal establishment (based on QS420, census 2011)
- **HOUSE_NRooms**: Number of rooms in the accommodation (LC4404EW, census 2011)
- **HOUSE_centralHeat**: Presence of central heating (based on LC4402EW, census 2011)
- **HOUSE_tenure**: Tenure (based on LC4402EW, census 2011)
- **HOUSE_NCars**: Number of cars (derived from LC4202EW by SPENSER team, census 2011)
- **id_HS**: unique identifier within the Health Survey (aggregated from the Health surveys from England, Wales and Scotland)
- **HEALTH_diabetes**: for Scotland and England, has doctor diagnosed diabetes; for Wales, diabetes currently treated (derived from HSE, HSW, SHS)
- **HEALTH_bloodpressure**: for Scotland and England, Doctor diagnosed high blood pressure; for Wales, high blood pressure currently treated (derived from HSE, HSW, SHS)
- **HEALTH_cvd**: for England, cardiovascular medication taken in the last 7 days; for Scotland, had cardiovascular condition excluding diabetes / blood pressure; for Wales, any heart condition excluding high blood pressure (derived from HSE, HSW, SHS)
- **HEALTH_NMedecines**: Number of prescribed medications (derived from HSE, HSW, SHS)
- **HEALTH_selfAssessed**: Self assessed general health (derived from HSE, HSW, SHS)
- **HEALTH_lifeSat**: how satisfied with life nowadays? (derived from HSE, HSW, SHS)
- **HEALTH_bmi**: BMI (see methods)
- **id_TUS_hh**: serial household identifier field in the UK Time Use Survey 2015
- **id_TUS_p**: pnun person identifier field in the UK Time Use Survey 2015
- **pwkstat**: Employment status (derived from UK TUS 2015)
- **soc2010**: Standard Occupational Classification (derived from UK TUS 2015)
- **sic1d2007**: Standard Industry Classification of economic activities 2007, 1st level (derived from UK TUS 2015)
- **sic2d2007**: Standard Industry Classification of economic activities 2007, 2nd level (derived from UK TUS 2015)
- **netPayWeekly**: Weekly take home pay after all deductions (derived from UK TUS 2015)
- **workedHoursWeekly**: Number of hours per week usually worked in main job or business (derived from UK TUS 2015)
- **incomeH**: Hourly gross salary for full-time and part-time employees (see methods)
- **incomeY**: Yearly gross salary for full-time and part-time employees (see methods)
- **incomeHAsIf**: Hourly gross salary for employees with self employed/other employees as employees of the same industry and with mean hourly worked for the industry when the number of hours is missing (see methods)
- **incomeYAsIf**: Yearly gross salary for employees with self employed/other employees as employees of the same industry and with mean hourly worked for the industry when the number of hours is missing (see methods)
- **ESport**: Relative probability weight to attend a sport fixture (Experimental, WIP)
- **ERugby**: Relative probability weight to attend a Rugby fixture (Experimental, WIP)
- **EConcertM**: Relative probability weight to attend a concert primarily targeting young males (Experimental, WIP)

- **EConcertF**: Relative probability weight to attend a concert primarily targeting young females (Experimental, WIP)
- **EConcertMS**: Relative probability weight to attend a concert primarily targeting middle-aged males (Experimental, WIP)
- **EConcertMS**: Relative probability weight to attend a concert primarily targeting middle-aged females (Experimental, WIP)
- **EMuseum**: Relative probability weight to visit a museum (Experimental, WIP)
- **easting**: X coordinate of the OA centroid in the British National Grid coordinate system (epsg:27700, source: ONS)
- **northing**: Y coordinate of the OA centroid in the British National Grid coordinate system (epsg:27700, source: ONS)
- **lng**: X coordinate of the OA centroid in the Longitude/Latitude coordinate system (epsg:4326, derived from ONS)
- **lat**: Y coordinate of the OA centroid in the Longitude/Latitude coordinate system (epsg:4326, derived from ONS)

13.3 National data

[businessRegistry.csv.gz](#)

Contains a breakdown of all business units (i.e. a single workplace) in Great Britain at LSOA scale, estimated by the project contributors from two nomis datasets: [UK Business Counts - local units by industry and employment size band 2020](#) and [Business Register and Employment Survey 2015](#). Each item contains the **size** of the unit and its main **sic1d07** code in reference to standard [Industrial Classification of Economic Activities 2007](#) (number corresponding to the letter in alphabetical order). It is used to compute commuting flows.

GIS/

Contains three GIS datasets of GB in GeoJson format taken from [ONS boundaries](#):

- [OA_2011_Pop20.geojson](#) at OA level
- [LSOA_2011_Pop20.geojson](#) at LSOA level
- [MSOA_2011_Pop20.geojson](#) at MSOA level

[QUANT_RAMP_spc.tar.gz](#)

See: Milton R, Batty M, Dennett A, dedicated [RAMP Spatial Interaction Model GitHub repository](#). It is used to compute the flows towards schools and retail.

[timeAtHomeIncreaseCTY.csv.gz](#)

This file is a subset from [Google COVID-19 Community Mobility Reports](#), cropped to GB. It describes the daily reduction in mobility, averaged at county level, due to lockdown and other COVID-19 restrictions between the 15th of February 2020 and 15th of October 2022. Missing values have been replaced by the national average. These values can be used directly to reduce **pnothome** and increase **phometot** (and their sub-categories) to simulate more accurately the period.

[diariesRef.csv.gz](#)

Contains diaries taken from the UK TUS that can be distributed to the population on a daily basis. They contain weekend days and weekday days. A full description of the fields can be found [here](#).

Part III

Advanced

14 Developer guide

14.1 Updating the docs

The site is built with [Quarto](#). You can iterate on it locally: `cd docs; quarto preview`

14.2 Code hygiene

We use automated tools to format the code.

```
cargo fmt

# Format Markdown docs
prettier --write *.md
prettier --write docs/*.qmd --parser markdown
```

Install [prettier](#) for Markdown.

14.3 Some tips for working with Rust

There are two equivalent ways to rebuild and then run the code. First:

```
cargo run --release -- devon
```

The `--` separates arguments to `cargo`, the Rust build tool, and arguments to the program itself. The second way:

```
cargo build --release
./target/release/aspics devon
```

You can build the code in two ways – **debug** and **release**. There's a simple tradeoff – debug mode is fast to build, but slow to run. Release mode is slow to build, but fast to run. For the ASPICS codebase, since the input data is so large and the codebase so small, I'd recommend always using `--release`. If you want to use debug mode, just omit the flag.

If you're working on the Rust code outside of an IDE like [VSCode](#), then you can check if the code compiles much faster by doing `cargo check`.

14.4 Docker

We provide a Dockerfile in case it's helpful for running, but don't recommend using it. If you want to, then assuming you have Docker setup:

```
docker build -t spc .
docker run --mount type=bind,source="$(pwd)"/data,target=/spc/data -t spc /spc/target/rele
```

This will make the `data` directory in your directory available to the Docker image, where it'll download the large input files and produce the final output.

15 Code walkthrough

SPC is implemented in [Rust](#), and its code can be found [here](#). This is an unusual implementation choice in the data science world, so this page has some notes about it.

15.1 Generally useful techniques

The code-base makes use of some techniques that may be generally applicable to other projects, independent of the language chosen.

15.1.1 Split code into two stages

Agent-based models and spatial interaction models require some kind of input. Often the effort to transform external data into this input can exceed that of the simulation component. Cleanly separating the two problems has some advantages:

- iterate on the simulation faster, without processing raw data every run
- reuse the prepared input for future projects
- force thinking about the data model needed by the simulation, and transform the external data into that form

SPC is exactly this first stage, originally split from [ASPICS](#) when further uses of the same population data were identified.

15.1.2 Explicit data schema

Dynamically typed languages like Python don't force you to explicitly list the shape of input data. It's common to read CSV files with `pandas`, filter and transform the data, and use that throughout the program. This can be quick to start prototyping, but is hard to maintain longer-term. Investing in the process of writing down types:

- makes it easier for somebody new to understand your system – they can first focus on **what** you're modeling, instead of how that's built up from raw data sources
- clarifies what data actually matters to your system; you don't carry forward unnecessary input

- makes it impossible to express invalid states
 - One example is [here](#) – per person and activity, there’s a list of venues the person may visit, along with a probability of going there. If the list of venues and list of probabilities are stored as separate lists or columns, then their length may not match.
- reuse the prepared input for future projects

There’s a variety of techniques for expressing strongly typed data:

- [protocol buffers](#) or [flatbuffers](#)
- [JSON schemas](#)
- [Python data classes](#) and [optional type hints](#)
- [statically typed languages like Rust](#)

15.1.3 Type-safe IDs

Say your data model has many different objects, each with their own ID – people, households, venues, etc. You might store these in a list and use the index as an ID. This is fine, but nothing stops you from confusing IDs and accidentally passing in venue 5 to a function instead of household 5. In Rust, it’s easy to create “wrapper types” like [this](#) and let the compiler prevent these mistakes.

This technique is also useful when preparing external data. [GTFS data](#) describing public transit routes and timetables contains many string IDs – shapes, trips, stops, routes. As soon as you read the raw input, you can [store the strings in more precise types](#) that prevent mixing up a stop ID and route ID.

15.1.4 Idempotent data preparation

If you’re iterating on your initialisation pipeline’s code, you probably don’t want to download a 2GB external file every single run. A common approach is to first test if a file exists and don’t download it again if so. In practice, you may also need to handle unzipping files, showing a progress bar while downloading, and printing clear error messages. This codebase has some [common code](#) for doing this in Rust. We intend to publish a separate library to more easily call in your own code.

15.1.5 Logging with structure

It's typical to print information as a complex pipeline runs, for the user to track progress and debug problems. But without any sort of organization, it's hard to follow what steps take a long time or encounter problems. What if your logs could show the logical structure of your pipeline and help you understand where time is spent?

```
[192.30s] [get_info_per_msoa] Loading buildings from data/raw_data/countydata/OSM/west-yorkshire-latest-free/
[192.64s] [get_info_per_msoa] Found 474,207 buildings from data/raw_data/countydata/OSM/west-yorkshire-latest-free/gis
osm_buildings_a_free_1.shp
[192.70s] [get_info_per_msoa] Matching 474,207 points to 299 polygons. Building R-Tree...
[194.22s] [calculate_lockdown_per_day] Calculating per-day lockdown values
[194.24s] [load_events] Loading events data
[194.25s] [initialisation] By the end, Memory usage: 1.53GiB
[200.89s] [Writing snapshot] Merging flows for all activities

212.24s      initialisation WestYorkshireLarge
31.18ms      grab_raw_data
192.04s      creating population
8.20s        read_individual_time_use_and_health_data
4.35s        Reading "data/raw_data/countydata/tus_hse_west-yorkshire.csv"
3.83s        Creating households
152.38s      create_commuting_flows
8.30s        setup_venue_flows Retail
6.59s        Copying flows to people Retail
7.47s        setup_venue_flows Nightclub
6.63s        Copying flows to people Nightclub
8.68s        setup_venue_flows PrimarySchool
6.58s        Copying flows to people PrimarySchool
7.00s        setup_venue_flows SecondarySchool
6.50s        Copying flows to people SecondarySchool
2.03s        get_info_per_msoa
24.48ms      calculate_lockdown_per_day
251.20µs     load_events "model_parameters/eventDataConcerts.csv"
1.07s        Writing population to "data/processed_data/WestYorkshireLarge/rust_cache.bin"
16.93s       Writing snapshot
```

The screenshot above shows a summary printed at the end of a long pipeline run. It's immediately obvious that the slowest step is creating commuting flows.

This codebase uses the [tracing](#) framework for logging, with a [custom piece](#) to draw the tree. (We'll publish this as a separate library once it's more polished.) The tracing framework is hard to understand, but the main conceptual leap over regular logging frameworks is the concept of a **span**. When your code starts one logical step, you call a method to create a new span, and when it finishes, you close that span. Spans can be nested in any way – `create_commuting_flows` happens within the larger step of `creating population`.

15.1.6 Determinism

Given the same inputs, your code should always produce identical output, no matter where it's run or how many times. Otherwise, debugging problems becomes very tedious, and it's more difficult to make conclusions from results. Of course, many projects have a stochastic element – but this should be controlled by a random number generator (RNG) seed, which is part of the input. You vary the seed and repeat the program, then reason about the distribution of results.

Aside from organizing your code to let a single RNG seed influence everything, another possible source of non-determinism is iteration order. In Rust, a `HashMap` could have different order every time it's used, so we use a `BTreeMap` instead when this matters. In Python, dictionaries are ordered. Be sure to check for your language.

15.2 Protocol buffers

SPC uses protocol buffers v2 for output. This has some advantages explained the “explicit data schema” section above.

Note that we chose proto2 instead of proto3, because proto3 doesn't support [required fields](#). This is done to allow schemas to evolve better over time, but this isn't a feature SPC makes use of. There's no need to have new code work with old data, or vice versa – if the schema is updated, downstream code should adapt accordingly and use the updated input files.

Note also that protocol buffers don't easily support type-safe wrappers around numeric IDs, so downstream code has to be careful not to mix up household, venue, and person IDs. For this reason, SPC internally doesn't use the auto-generated protobuf code until the very end of the pipeline. It's always possible to be more precise with native Rust types, and convert to the less strict types later.

15.3 An example of the power of static type checking

Imagine we want to add a new activity type to represent people going to university and higher education. SPC already has activities for primary and secondary school, so we'll probably want to follow those as a guide. In any language, we could search the codebase for relevant terms to get a sense of what to update. In languages like Python without an up-front compilation step, if we fail to update something or write blatantly incorrect code (such as making a typo in variable names or passing a list where a string was expected), we only find out when that code happens to run. In pipelines with many steps and large input files, it could be a while before we reach the problematic code.

Let's walk through the same exercise for SPC's Rust code. We start by adding a new `University` case to the [Activity enum](#). If we try to compile the code here (with `cargo check` or an IDE), we immediately get 4 errors.

```

error[E0004]: non-exhaustive patterns: `University` not covered
--> src/init/quant.rs:38:44
38 |         let (population_csv, prob_sij) = match activity {
    |                                           ^^^^^^^^^ pattern `University` not covered
    |
    = help: ensure that all possible cases are being handled, possibly by adding wildcards or more
match arms
    = note: the matched value is of type `Activity`
    ::: src/lib.rs:129:1
129 | / pub enum Activity {
130 | |     Retail,
131 | |     PrimarySchool,
132 | |     SecondarySchool,
    | |     ...
135 | |     University,
    | |     ----- not covered
136 | | }
    | |_- `Activity` defined here

```

Three of the errors are in the QUANT module. The first is [here](#). It's immediately clear that for retail and primary/secondary school, we read in two files from QUANT representing venues where these activities take place and the probability of going to each venue. Even if we were unfamiliar with this codebase, the compiler has told us one thing we'll need to figure out, and where to wire it up.

```

error[E0004]: non-exhaustive patterns: `University` not covered
--> src/protobuf.rs:135:11
135 |         match activity {
    |                   ^^^^^^^^^ pattern `University` not covered
    |
    = help: ensure that all possible cases are being handled, possibly by adding wildcards or more
match arms
    = note: the matched value is of type `Activity`
    ::: src/lib.rs:129:1
129 | / pub enum Activity {
130 | |     Retail,
131 | |     PrimarySchool,
132 | |     SecondarySchool,
    | |     ...
135 | |     University,
    | |     ----- not covered
136 | | }
    | |_- `Activity` defined here

```

The other error is in the [code that writes the protobuf output](#). Similarly, we need a way to represent university activities in the protobuf scheme.

Extending an unfamiliar code-base backed by compiler errors is a very guided experience. If you wanted to add more demographic attributes to people or energy use information to households, you don't need to guess all of the places in the code you'll need to update. You can just add the field, then let the compiler tell you all places where those objects get created.

16 Performance

The following tables summarizes the resources SPC needs to run in different areas.

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|----------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2012 | England/bedfordshire | 74 | 245,166 | 647,272 | 256.83 MiB | 10 seconds | 3 seconds | 849.02 MiB |
| 2020 | England/bedfordshire | 74 | 272,875 | 674,044 | 271.65 MiB | 9 seconds | 3 seconds | 922.88 MiB |
| 2022 | England/bedfordshire | 74 | 309,706 | 703,582 | 277.74 MiB | 9 seconds | 3 seconds | 929.80 MiB |
| 2032 | England/bedfordshire | 74 | 309,706 | 703,582 | 277.74 MiB | 9 seconds | 3 seconds | 929.80 MiB |
| 2039 | England/bedfordshire | 74 | 329,061 | 715,797 | 278.39 MiB | 11 seconds | 3 seconds | 927.77 MiB |
| 2012 | England/berkshire | 107 | 342,167 | 890,543 | 356.04 MiB | 14 seconds | 7 seconds | 1.06 GiB |
| 2020 | England/berkshire | 107 | 365,905 | 918,258 | 373.35 MiB | 14 seconds | 7 seconds | 1.10 GiB |
| 2022 | England/berkshire | 107 | 394,446 | 941,655 | 368.37 MiB | 14 seconds | 7 seconds | 1.08 GiB |
| 2032 | England/berkshire | 107 | 394,446 | 941,655 | 368.37 MiB | 14 seconds | 7 seconds | 1.08 GiB |
| 2039 | England/berkshire | 107 | 408,604 | 949,986 | 367.21 MiB | 14 seconds | 7 seconds | 1.08 GiB |
| 2012 | England/bristol | 55 | 182,299 | 448,233 | 173.74 MiB | 6 seconds | 2 seconds | 527.23 MiB |
| 2020 | England/bristol | 55 | 196,940 | 470,039 | 183.99 MiB | 6 seconds | 2 seconds | 547.49 MiB |

| year | study_area | num_msas | soas_households | schoolspeople | file_size | runtime | commuting_time | memory_usage |
|------|-----------------------|----------|-----------------|---------------|------------|------------|----------------|--------------|
| 2022 | England/bristol | 55 | 216,197 | 503,014 | 192.51 MiB | 7 seconds | 2 seconds | 559.78 MiB |
| 2032 | England/bristol | 55 | 216,197 | 503,014 | 192.51 MiB | 7 seconds | 2 seconds | 559.78 MiB |
| 2039 | England/bristol | 55 | 227,770 | 521,371 | 199.72 MiB | 7 seconds | 2 seconds | 573.40 MiB |
| 2012 | England/buckinghamsh | 99 | 99,235 | 261,340 | 108.30 MiB | 5 seconds | 1 second | 310.38 MiB |
| 2020 | England/buckinghamsh | 99 | 108,999 | 271,050 | 114.31 MiB | 5 seconds | 1 second | 400.87 MiB |
| 2022 | England/buckinghamsh | 99 | 123,578 | 278,548 | 112.39 MiB | 5 seconds | 1 second | 393.64 MiB |
| 2032 | England/buckinghamsh | 99 | 123,578 | 278,548 | 112.39 MiB | 5 seconds | 1 second | 393.64 MiB |
| 2039 | England/buckinghamsh | 99 | 130,393 | 281,773 | 112.14 MiB | 5 seconds | 1 second | 391.52 MiB |
| 2012 | England/cambridgeshir | 98 | 327,257 | 832,980 | 323.35 MiB | 11 seconds | 5 seconds | 1013.16 MiB |
| 2020 | England/cambridgeshir | 98 | 348,522 | 863,250 | 341.16 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2022 | England/cambridgeshir | 98 | 377,634 | 907,166 | 348.75 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2032 | England/cambridgeshir | 98 | 377,634 | 907,166 | 348.75 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2039 | England/cambridgeshir | 98 | 392,478 | 924,170 | 351.39 MiB | 12 seconds | 5 seconds | 1.04 GiB |
| 2012 | England/cheshire | 139 | 441,084 | 1,042,064 | 402.14 MiB | 16 seconds | 7 seconds | 1.13 GiB |
| 2020 | England/cheshire | 139 | 464,134 | 1,070,597 | 416.35 MiB | 16 seconds | 7 seconds | 1.46 GiB |
| 2022 | England/cheshire | 139 | 489,476 | 1,125,198 | 425.27 MiB | 16 seconds | 7 seconds | 1.47 GiB |

| year | study_area | num_msas | houses | households | people | file_size | runtime | commuting | memory_usage |
|------|--------------------|----------|---------|------------|---------------|--------------------|-----------|---------------|--------------|
| 2032 | England/cheshire | 139 | 489,476 | 1,125,198 | 125.27 MiB | 16 sec- onds | 7 seconds | 1.47 GiB | |
| 2039 | England/cheshire | 139 | 501,501 | 1,149,514 | 131.10 MiB | 16 sec- onds | 7 seconds | 1.48 GiB | |
| 2012 | England/cornwall | 74 | 232,659 | 549,616 | 208.07 MiB | 8 sec- onds | 2 seconds | 742.92 MiB | |
| 2020 | England/cornwall | 74 | 247,105 | 577,414 | 219.74 MiB | 8 sec- onds | 3 seconds | 764.95 MiB | |
| 2022 | England/cornwall | 74 | 270,134 | 634,940 | 233.36 MiB | 8 sec- onds | 3 seconds | 828.45 MiB | |
| 2032 | England/cornwall | 74 | 270,134 | 634,940 | 233.36 MiB | 8 sec- onds | 3 seconds | 828.45 MiB | |
| 2039 | England/cornwall | 74 | 280,546 | 658,610 | 239.72 MiB | 8 sec- onds | 3 seconds | 838.14 MiB | |
| 2012 | England/cumbria | 64 | 222,586 | 498,624 | 188.03 MiB | 7 sec- onds | 2 seconds | 547.33 MiB | |
| 2020 | England/cumbria | 64 | 226,893 | 499,873 | 188.73 MiB | 7 sec- onds | 2 seconds | 548.52 MiB | |
| 2022 | England/cumbria | 64 | 230,206 | 499,840 | 183.18 MiB | 7 sec- onds | 2 seconds | 533.99 MiB | |
| 2032 | England/cumbria | 64 | 230,206 | 499,840 | 183.18 MiB | 7 sec- onds | 2 seconds | 533.99 MiB | |
| 2039 | England/cumbria | 64 | 231,202 | 498,475 | 181.58 MiB | 7 sec- onds | 2 seconds | 530.96 MiB | |
| 2012 | England/derbyshire | 131 | 436,276 | 1,035,356 | 197.75 MiB | 15 sec- onds | 7 seconds | 1.12 GiB | |
| 2020 | England/derbyshire | 131 | 459,743 | 1,064,406 | 109.76 MiB | 16 sec- onds | 8 seconds | 1.44 GiB | |
| 2022 | England/derbyshire | 131 | 489,764 | 1,122,078 | 119.52 MiB | 16 sec- onds | 7 seconds | 1.45 GiB | |
| 2032 | England/derbyshire | 131 | 489,764 | 1,122,078 | 119.52 MiB | 16 sec- onds | 7 seconds | 1.45 GiB | |
| 2039 | England/derbyshire | 131 | 505,314 | 1,152,518 | 129.01 MiB | 16 sec- onds | 8 seconds | 1.47 GiB | |

| year | study_area | num_msaas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|---------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2012 | England/devon | 156 | 494,106 | 1,165,952 | 238.60 MiB | 18 seconds | 8 seconds | 1.49 GiB |
| 2020 | England/devon | 156 | 523,033 | 1,212,387 | 259.44 MiB | 18 seconds | 8 seconds | 1.53 GiB |
| 2022 | England/devon | 156 | 567,011 | 1,304,874 | 278.71 MiB | 19 seconds | 9 seconds | 1.64 GiB |
| 2032 | England/devon | 156 | 567,011 | 1,304,874 | 278.71 MiB | 19 seconds | 9 seconds | 1.64 GiB |
| 2039 | England/devon | 156 | 589,178 | 1,342,774 | 288.23 MiB | 19 seconds | 9 seconds | 1.66 GiB |
| 2012 | England/durham | 117 | 390,472 | 911,601 | 349.78 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2020 | England/durham | 117 | 407,828 | 930,184 | 359.59 MiB | 12 seconds | 5 seconds | 1.05 GiB |
| 2022 | England/durham | 117 | 425,611 | 952,801 | 356.63 MiB | 12 seconds | 5 seconds | 1.04 GiB |
| 2032 | England/durham | 117 | 425,611 | 952,801 | 356.63 MiB | 12 seconds | 5 seconds | 1.04 GiB |
| 2039 | England/durham | 117 | 434,593 | 959,555 | 357.66 MiB | 12 seconds | 5 seconds | 1.04 GiB |
| 2012 | England/east-sussex | 102 | 355,257 | 827,703 | 313.71 MiB | 12 seconds | 5 seconds | 987.31 MiB |
| 2020 | England/east-sussex | 102 | 380,894 | 853,970 | 324.02 MiB | 12 seconds | 6 seconds | 1006.13 MiB |
| 2022 | England/east-sussex | 102 | 423,181 | 895,907 | 329.56 MiB | 12 seconds | 5 seconds | 1008.58 MiB |

| year | study_area | num_msoas | soas_households | people | file_size | runtime | commuting_time | memory_usage |
|------|----------------------------------|-----------|-----------------|-----------|------------|------------|----------------|--------------|
| 2032 | England/east-sussex | 102 | 423,181 | 895,907 | 329.56 MiB | 12 seconds | 5 seconds | 1008.58 MiB |
| 2039 | England/east-sussex | 102 | 446,000 | 915,014 | 335.45 MiB | 12 seconds | 5 seconds | 1020.75 MiB |
| 2012 | England/east-yorkshire-with-hull | 75 | 255,848 | 593,271 | 227.49 MiB | 8 seconds | 3 seconds | 778.75 MiB |
| 2020 | England/east-yorkshire-with-hull | 75 | 262,609 | 602,286 | 233.14 MiB | 8 seconds | 3 seconds | 835.04 MiB |
| 2022 | England/east-yorkshire-with-hull | 75 | 272,805 | 613,721 | 230.34 MiB | 8 seconds | 3 seconds | 824.50 MiB |
| 2032 | England/east-yorkshire-with-hull | 75 | 272,805 | 613,721 | 230.34 MiB | 8 seconds | 3 seconds | 824.50 MiB |
| 2039 | England/east-yorkshire-with-hull | 75 | 277,770 | 617,357 | 230.45 MiB | 8 seconds | 3 seconds | 825.00 MiB |
| 2012 | England/essex | 211 | 722,974 | 1,786,316 | 690.77 MiB | 30 seconds | 19 seconds | 2.06 GiB |
| 2020 | England/essex | 211 | 773,454 | 1,857,205 | 726.02 MiB | 32 seconds | 20 seconds | 2.13 GiB |
| 2022 | England/essex | 211 | 858,552 | 1,981,994 | 761.40 MiB | 33 seconds | 20 seconds | 2.19 GiB |
| 2032 | England/essex | 211 | 858,552 | 1,981,994 | 761.40 MiB | 33 seconds | 20 seconds | 2.19 GiB |
| 2039 | England/essex | 211 | 906,640 | 2,042,404 | 777.71 MiB | 33 seconds | 21 seconds | 2.21 GiB |
| 2012 | England/gloucestershire107 | 107 | 365,240 | 889,836 | 344.16 MiB | 13 seconds | 5 seconds | 1.02 GiB |
| 2020 | England/gloucestershire107 | 107 | 392,643 | 933,909 | 362.90 MiB | 13 seconds | 6 seconds | 1.06 GiB |
| 2022 | England/gloucestershire107 | 107 | 432,216 | 1,025,073 | 389.56 MiB | 14 seconds | 6 seconds | 1.10 GiB |

| year | study_area | num_msoas | num_households | num_people | file_size | time commuting | memory usage |
|------|----------------------------|-----------|----------------|------------|------------|----------------|--------------|
| 2032 | England/gloucestershire | 107 | 432,216 | 1,025,073 | 89.56 MiB | 14 seconds | 1.10 GiB |
| 2039 | England/gloucestershire | 107 | 453,383 | 1,068,484 | 103.87 MiB | 14 seconds | 1.43 GiB |
| 2012 | England/greater-london | 983 | 3,283,305 | 8,581,243 | 3.27 GiB | 12 minutes | 11.80 GiB |
| 2020 | England/greater-london | 983 | 3,574,266 | 8,983,773 | 3.48 GiB | 12 minutes | 12.22 GiB |
| 2022 | England/greater-london | 983 | 3,997,548 | 9,452,043 | 3.55 GiB | 12 minutes | 12.25 GiB |
| 2032 | England/greater-london | 983 | 3,997,548 | 9,452,043 | 3.55 GiB | 12 minutes | 12.25 GiB |
| 2039 | England/greater-london | 983 | 4,229,017 | 9,688,503 | 3.58 GiB | 12 minutes | 12.95 GiB |
| 2012 | England/greater-manchester | 346 | 1,128,371 | 2,745,451 | 1.05 GiB | 70 seconds | 3.56 GiB |
| 2020 | England/greater-manchester | 346 | 1,192,547 | 2,840,431 | 1.10 GiB | 71 seconds | 3.66 GiB |
| 2022 | England/greater-manchester | 346 | 1,272,689 | 2,974,954 | 1.13 GiB | 73 seconds | 3.69 GiB |
| 2032 | England/greater-manchester | 346 | 1,272,689 | 2,974,954 | 1.13 GiB | 74 seconds | 3.69 GiB |
| 2039 | England/greater-manchester | 346 | 1,319,090 | 3,049,727 | 1.15 GiB | 76 seconds | 3.73 GiB |
| 2012 | England/hampshire | 225 | 733,611 | 1,810,518 | 98.03 MiB | 32 seconds | 2.07 GiB |

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|-----------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2020 | England/hampshire | 225 | 777,116 | 1,861,250 | 721.62 MiB | 32 seconds | 20 seconds | 2.12 GiB |
| 2022 | England/hampshire | 225 | 836,451 | 1,931,667 | 728.97 MiB | 32 seconds | 20 seconds | 2.12 GiB |
| 2032 | England/hampshire | 225 | 836,451 | 1,931,667 | 728.97 MiB | 33 seconds | 20 seconds | 2.12 GiB |
| 2039 | England/hampshire | 225 | 867,417 | 1,960,190 | 735.50 MiB | 33 seconds | 20 seconds | 2.13 GiB |
| 2012 | England/herefordshire | 23 | 79,083 | 188,362 | 72.21 MiB | 4 seconds | 1 second | 234.89 MiB |
| 2020 | England/herefordshire | 23 | 83,238 | 195,194 | 74.71 MiB | 4 seconds | 1 second | 239.36 MiB |
| 2022 | England/herefordshire | 23 | 89,574 | 209,784 | 77.63 MiB | 4 seconds | 1 second | 242.83 MiB |
| 2032 | England/herefordshire | 23 | 89,574 | 209,784 | 77.63 MiB | 4 seconds | 1 second | 242.83 MiB |
| 2039 | England/herefordshire | 23 | 92,605 | 216,508 | 79.43 MiB | 4 seconds | 1 second | 245.69 MiB |
| 2012 | England/hertfordshire | 153 | 457,276 | 1,160,151 | 458.65 MiB | 19 seconds | 11 seconds | 1.56 GiB |
| 2020 | England/hertfordshire | 153 | 494,661 | 1,190,041 | 477.18 MiB | 19 seconds | 11 seconds | 1.59 GiB |
| 2022 | England/hertfordshire | 153 | 546,573 | 1,219,124 | 476.55 MiB | 19 seconds | 11 seconds | 1.67 GiB |
| 2032 | England/hertfordshire | 153 | 546,573 | 1,219,124 | 476.55 MiB | 19 seconds | 11 seconds | 1.67 GiB |
| 2039 | England/hertfordshire | 153 | 575,179 | 1,233,571 | 476.97 MiB | 19 seconds | 11 seconds | 1.67 GiB |
| 2012 | England/isle-of-wight | 18 | 61,636 | 139,732 | 53.88 MiB | 3 seconds | 1 second | 188.79 MiB |
| 2020 | England/isle-of-wight | 18 | 65,140 | 143,268 | 54.99 MiB | 3 seconds | 1 second | 190.45 MiB |

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|------------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2022 | England/isle-of-wight | 18 | 70,496 | 151,582 | 55.55 MiB | 3 seconds | 1 second | 200.99 MiB |
| 2032 | England/isle-of-wight | 18 | 70,496 | 151,582 | 55.55 MiB | 3 seconds | 1 second | 200.99 MiB |
| 2039 | England/isle-of-wight | 18 | 72,968 | 154,841 | 56.14 MiB | 3 seconds | 1 second | 202.13 MiB |
| 2012 | England/kent | 220 | 718,544 | 1,793,707 | 200.10 MiB | 29 seconds | 17 seconds | 2.08 GiB |
| 2020 | England/kent | 220 | 781,933 | 1,873,457 | 237.20 MiB | 30 seconds | 18 seconds | 2.15 GiB |
| 2022 | England/kent | 220 | 875,515 | 2,008,857 | 273.24 MiB | 31 seconds | 19 seconds | 2.21 GiB |
| 2032 | England/kent | 220 | 875,515 | 2,008,857 | 273.24 MiB | 32 seconds | 19 seconds | 2.21 GiB |
| 2039 | England/kent | 220 | 926,571 | 2,069,087 | 288.47 MiB | 35 seconds | 19 seconds | 2.23 GiB |
| 2012 | England/lancashire | 191 | 619,861 | 1,476,465 | 171.94 MiB | 24 seconds | 14 seconds | 1.83 GiB |
| 2020 | England/lancashire | 191 | 640,196 | 1,511,896 | 189.78 MiB | 24 seconds | 14 seconds | 1.87 GiB |
| 2022 | England/lancashire | 191 | 663,637 | 1,567,396 | 194.49 MiB | 24 seconds | 14 seconds | 1.87 GiB |
| 2032 | England/lancashire | 191 | 663,637 | 1,567,396 | 194.49 MiB | 24 seconds | 14 seconds | 1.87 GiB |
| 2039 | England/lancashire | 191 | 674,387 | 1,591,908 | 200.02 MiB | 25 seconds | 14 seconds | 1.88 GiB |
| 2012 | England/leicestershire | 120 | 370,305 | 958,470 | 373.02 MiB | 14 seconds | 7 seconds | 1.08 GiB |

| year | study_area | num_msoas | households | people | file_size | runtime | commuting_time | memory_usage |
|------|------------------------|-----------|------------|-----------|------------|------------|----------------|--------------|
| 2020 | England/leicestershire | 120 | 397,467 | 1,016,632 | 297.28 MiB | 14 seconds | 7 seconds | 1.13 GiB |
| 2022 | England/leicestershire | 120 | 438,413 | 1,118,737 | 426.12 MiB | 15 seconds | 8 seconds | 1.48 GiB |
| 2032 | England/leicestershire | 120 | 438,413 | 1,118,737 | 426.12 MiB | 15 seconds | 8 seconds | 1.48 GiB |
| 2039 | England/leicestershire | 120 | 459,655 | 1,164,678 | 440.87 MiB | 16 seconds | 8 seconds | 1.50 GiB |
| 2012 | England/lincolnshire | 134 | 449,394 | 1,064,403 | 403.05 MiB | 15 seconds | 7 seconds | 1.43 GiB |
| 2020 | England/lincolnshire | 134 | 475,646 | 1,098,403 | 419.31 MiB | 15 seconds | 7 seconds | 1.46 GiB |
| 2022 | England/lincolnshire | 134 | 507,295 | 1,152,299 | 427.55 MiB | 15 seconds | 7 seconds | 1.47 GiB |
| 2032 | England/lincolnshire | 134 | 507,295 | 1,152,299 | 427.55 MiB | 16 seconds | 7 seconds | 1.47 GiB |
| 2039 | England/lincolnshire | 134 | 523,548 | 1,172,923 | 430.83 MiB | 16 seconds | 7 seconds | 1.47 GiB |
| 2012 | England/merseyside | 184 | 603,483 | 1,399,205 | 533.96 MiB | 20 seconds | 11 seconds | 1.75 GiB |
| 2020 | England/merseyside | 184 | 632,617 | 1,435,755 | 553.33 MiB | 20 seconds | 11 seconds | 1.79 GiB |
| 2022 | England/merseyside | 184 | 665,766 | 1,498,518 | 570.21 MiB | 20 seconds | 10 seconds | 1.82 GiB |
| 2032 | England/merseyside | 184 | 665,766 | 1,498,518 | 570.21 MiB | 21 seconds | 11 seconds | 1.82 GiB |

| year | study_area | num_msoas | soas_households | householdspeople | file_size | runtime | commuting_time | memory_usage |
|------|--------------------------|-----------|-----------------|------------------|------------|------------|----------------|--------------|
| 2039 | England/merseyside | 184 | 685,165 | 1,528,037 | 777.48 MiB | 21 seconds | 11 seconds | 1.83 GiB |
| 2012 | England/norfolk | 110 | 374,491 | 882,793 | 333.07 MiB | 12 seconds | 5 seconds | 1017.16 MiB |
| 2020 | England/norfolk | 110 | 397,770 | 916,799 | 348.41 MiB | 12 seconds | 5 seconds | 1.02 GiB |
| 2022 | England/norfolk | 110 | 432,187 | 982,755 | 362.27 MiB | 13 seconds | 5 seconds | 1.04 GiB |
| 2032 | England/norfolk | 110 | 432,187 | 982,755 | 362.27 MiB | 13 seconds | 5 seconds | 1.04 GiB |
| 2039 | England/norfolk | 110 | 450,068 | 1,013,214 | 471.39 MiB | 13 seconds | 5 seconds | 1.06 GiB |
| 2012 | England/northamptonshire | 91 | 289,575 | 720,263 | 284.39 MiB | 10 seconds | 4 seconds | 941.33 MiB |
| 2020 | England/northamptonshire | 91 | 316,553 | 762,382 | 304.36 MiB | 10 seconds | 4 seconds | 981.14 MiB |
| 2022 | England/northamptonshire | 91 | 352,529 | 828,003 | 320.81 MiB | 11 seconds | 5 seconds | 1005.64 MiB |
| 2032 | England/northamptonshire | 91 | 352,529 | 828,003 | 320.81 MiB | 11 seconds | 5 seconds | 1005.64 MiB |
| 2039 | England/northamptonshire | 91 | 370,555 | 855,812 | 328.03 MiB | 11 seconds | 5 seconds | 1016.86 MiB |
| 2012 | England/northumberland | 40 | 138,928 | 315,894 | 120.66 MiB | 5 seconds | 1 second | 423.11 MiB |
| 2020 | England/northumberland | 40 | 143,516 | 322,616 | 121.94 MiB | 5 seconds | 1 second | 423.87 MiB |
| 2022 | England/northumberland | 40 | 148,792 | 333,456 | 122.06 MiB | 5 seconds | 1 second | 421.48 MiB |
| 2032 | England/northumberland | 40 | 148,792 | 333,456 | 122.06 MiB | 5 seconds | 1 second | 421.48 MiB |

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|-------------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2039 | England/northumberland | 40 | 150,259 | 337,186 | 122.24 MiB | 5 seconds | 1 second | 421.48 MiB |
| 2012 | England/north-yorkshire | 138 | 460,050 | 1,085,067 | 113.05 MiB | 16 seconds | 7 seconds | 1.45 GiB |
| 2020 | England/north-yorkshire | 138 | 478,639 | 1,107,928 | 123.18 MiB | 16 seconds | 7 seconds | 1.47 GiB |
| 2022 | England/north-yorkshire | 138 | 499,392 | 1,134,723 | 120.60 MiB | 16 seconds | 7 seconds | 1.45 GiB |
| 2032 | England/north-yorkshire | 138 | 499,392 | 1,134,723 | 120.60 MiB | 16 seconds | 7 seconds | 1.45 GiB |
| 2039 | England/north-yorkshire | 138 | 509,099 | 1,143,897 | 121.52 MiB | 16 seconds | 7 seconds | 1.46 GiB |
| 2012 | England/nottinghamshire | 138 | 460,022 | 1,123,007 | 132.35 MiB | 16 seconds | 8 seconds | 1.49 GiB |
| 2020 | England/nottinghamshire | 138 | 486,163 | 1,169,489 | 153.68 MiB | 16 seconds | 8 seconds | 1.53 GiB |
| 2022 | England/nottinghamshire | 138 | 522,944 | 1,248,804 | 173.35 MiB | 17 seconds | 8 seconds | 1.56 GiB |
| 2032 | England/nottinghamshire | 138 | 522,944 | 1,248,804 | 173.35 MiB | 17 seconds | 8 seconds | 1.56 GiB |
| 2039 | England/nottinghamshire | 138 | 543,291 | 1,281,812 | 182.21 MiB | 18 seconds | 9 seconds | 1.66 GiB |
| 2012 | England/oxfordshire | 86 | 261,235 | 671,997 | 260.43 MiB | 9 seconds | 3 seconds | 852.84 MiB |
| 2020 | England/oxfordshire | 86 | 274,908 | 695,490 | 271.62 MiB | 10 seconds | 4 seconds | 918.90 MiB |
| 2022 | England/oxfordshire | 86 | 293,368 | 729,866 | 275.40 MiB | 10 seconds | 4 seconds | 919.34 MiB |

| year | study_area | num_msas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|---------------------|----------|----------------|------------|------------|------------|----------------|--------------|
| 2032 | England/oxfordshire | 86 | 293,368 | 729,866 | 275.40 MiB | 10 seconds | 4 seconds | 919.34 MiB |
| 2039 | England/oxfordshire | 86 | 303,035 | 743,227 | 277.51 MiB | 10 seconds | 4 seconds | 922.19 MiB |
| 2012 | England/rutland | 5 | 14,912 | 38,314 | 16.37 MiB | 3 seconds | 1 second | 54.07 MiB |
| 2020 | England/rutland | 5 | 16,698 | 40,381 | 17.09 MiB | 3 seconds | 1 second | 57.95 MiB |
| 2022 | England/rutland | 5 | 18,198 | 44,193 | 18.26 MiB | 3 seconds | 1 second | 60.08 MiB |
| 2032 | England/rutland | 5 | 18,198 | 44,193 | 18.26 MiB | 3 seconds | 1 second | 60.08 MiB |
| 2039 | England/rutland | 5 | 18,914 | 45,659 | 18.71 MiB | 3 seconds | 1 second | 61.20 MiB |
| 2012 | England/shropshire | 62 | 197,768 | 483,414 | 186.29 MiB | 7 seconds | 2 seconds | 550.97 MiB |
| 2020 | England/shropshire | 62 | 211,035 | 508,233 | 195.76 MiB | 7 seconds | 2 seconds | 568.62 MiB |
| 2022 | England/shropshire | 62 | 228,285 | 558,755 | 207.29 MiB | 7 seconds | 2 seconds | 740.58 MiB |
| 2032 | England/shropshire | 62 | 228,285 | 558,755 | 207.29 MiB | 7 seconds | 2 seconds | 740.58 MiB |
| 2039 | England/shropshire | 62 | 236,015 | 581,476 | 213.23 MiB | 7 seconds | 2 seconds | 749.82 MiB |
| 2012 | England/somerset | 124 | 329,040 | 790,346 | 303.62 MiB | 11 seconds | 4 seconds | 970.03 MiB |
| 2020 | England/somerset | 124 | 353,976 | 822,271 | 317.43 MiB | 11 seconds | 4 seconds | 996.71 MiB |
| 2022 | England/somerset | 124 | 388,675 | 880,441 | 331.61 MiB | 12 seconds | 4 seconds | 1018.53 MiB |
| 2032 | England/somerset | 124 | 388,675 | 880,441 | 331.61 MiB | 12 seconds | 4 seconds | 1018.53 MiB |
| 2039 | England/somerset | 124 | 406,157 | 906,545 | 339.87 MiB | 12 seconds | 4 seconds | 1.01 GiB |

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|-------------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2012 | England/south-yorkshire | 172 | 566,664 | 1,372,435 | 28.11 MiB | 20 seconds | 11 seconds | 1.75 GiB |
| 2020 | England/south-yorkshire | 172 | 597,694 | 1,418,846 | 48.59 MiB | 21 seconds | 11 seconds | 1.79 GiB |
| 2032 | England/south-yorkshire | 172 | 637,411 | 1,493,545 | 63.91 MiB | 21 seconds | 11 seconds | 1.81 GiB |
| 2039 | England/south-yorkshire | 172 | 659,843 | 1,531,313 | 75.31 MiB | 22 seconds | 12 seconds | 1.83 GiB |
| 2012 | England/staffordshire | 143 | 464,441 | 1,111,144 | 25.27 MiB | 16 seconds | 8 seconds | 1.47 GiB |
| 2020 | England/staffordshire | 143 | 486,645 | 1,139,752 | 37.51 MiB | 16 seconds | 8 seconds | 1.49 GiB |
| 2022 | England/staffordshire | 143 | 510,634 | 1,188,857 | 44.87 MiB | 17 seconds | 8 seconds | 1.50 GiB |
| 2032 | England/staffordshire | 143 | 510,634 | 1,188,857 | 44.87 MiB | 17 seconds | 8 seconds | 1.50 GiB |
| 2039 | England/staffordshire | 143 | 522,882 | 1,215,006 | 52.94 MiB | 17 seconds | 8 seconds | 1.52 GiB |
| 2012 | England/suffolk | 90 | 136,142 | 327,349 | 128.13 MiB | 5 seconds | 1 second | 440.37 MiB |
| 2020 | England/suffolk | 90 | 146,277 | 333,781 | 130.90 MiB | 5 seconds | 1 second | 445.14 MiB |
| 2022 | England/suffolk | 90 | 159,882 | 344,534 | 130.76 MiB | 5 seconds | 1 second | 442.14 MiB |
| 2032 | England/suffolk | 90 | 159,882 | 344,534 | 130.76 MiB | 5 seconds | 1 second | 442.14 MiB |
| 2039 | England/suffolk | 90 | 166,718 | 350,358 | 132.54 MiB | 5 seconds | 1 second | 446.24 MiB |
| 2012 | England/surrey | 151 | 458,108 | 1,168,112 | 56.50 MiB | 21 seconds | 13 seconds | 1.55 GiB |

| year | study_area | num_msaas | households | people | file_size | runtime | commuting_time | memory_usage |
|------|-----------------------|-----------|------------|-----------|------------|------------|----------------|--------------|
| 2020 | England/surrey | 151 | 480,930 | 1,195,509 | 172.89 MiB | 21 seconds | 13 seconds | 1.58 GiB |
| 2022 | England/surrey | 151 | 518,720 | 1,214,557 | 167.03 MiB | 21 seconds | 13 seconds | 1.56 GiB |
| 2032 | England/surrey | 151 | 518,720 | 1,214,557 | 167.03 MiB | 21 seconds | 13 seconds | 1.56 GiB |
| 2039 | England/surrey | 151 | 538,941 | 1,221,227 | 164.71 MiB | 21 seconds | 13 seconds | 1.64 GiB |
| 2012 | England/tyne-and-wear | 145 | 483,909 | 1,119,030 | 127.35 MiB | 15 seconds | 7 seconds | 1.47 GiB |
| 2020 | England/tyne-and-wear | 145 | 501,383 | 1,143,194 | 139.09 MiB | 15 seconds | 7 seconds | 1.50 GiB |
| 2022 | England/tyne-and-wear | 145 | 521,777 | 1,168,078 | 140.03 MiB | 15 seconds | 7 seconds | 1.49 GiB |
| 2032 | England/tyne-and-wear | 145 | 521,777 | 1,168,078 | 140.03 MiB | 14 seconds | 6 seconds | 1.49 GiB |
| 2039 | England/tyne-and-wear | 145 | 532,652 | 1,177,340 | 141.36 MiB | 15 seconds | 7 seconds | 1.58 GiB |
| 2012 | England/warwickshire | 108 | 361,467 | 896,673 | 347.44 MiB | 13 seconds | 6 seconds | 1.03 GiB |
| 2020 | England/warwickshire | 108 | 392,639 | 958,833 | 373.63 MiB | 14 seconds | 6 seconds | 1.08 GiB |
| 2022 | England/warwickshire | 108 | 432,682 | 1,061,957 | 405.95 MiB | 15 seconds | 7 seconds | 1.44 GiB |
| 2032 | England/warwickshire | 108 | 432,682 | 1,061,957 | 405.95 MiB | 14 seconds | 7 seconds | 1.44 GiB |

| year | study_area | num_msoas | num_households | num_people | file_size | runtime | commuting_time | memory_usage |
|------|------------------------|-----------|----------------|------------|------------|------------|----------------|--------------|
| 2039 | England/warwickshire | 108 | 454,732 | 1,112,230 | 124.10 MiB | 15 seconds | 7 seconds | 1.47 GiB |
| 2012 | England/west-midlands | 314 | 958,034 | 2,477,399 | 190.27 MiB | 56 seconds | 38 seconds | 3.24 GiB |
| 2020 | England/west-midlands | 314 | 1,002,273 | 2,572,395 | 1.01 GiB | 58 seconds | 40 seconds | 3.33 GiB |
| 2022 | England/west-midlands | 314 | 1,046,146 | 2,664,228 | 1.04 GiB | 60 seconds | 41 seconds | 3.37 GiB |
| 2032 | England/west-midlands | 314 | 1,079,612 | 2,706,242 | 1.04 GiB | 61 seconds | 41 seconds | 3.55 GiB |
| 2039 | England/west-midlands | 314 | 1,128,890 | 2,787,990 | 1.07 GiB | 62 seconds | 42 seconds | 3.59 GiB |
| 2012 | England/west-sussex | 100 | 348,766 | 836,646 | 321.17 MiB | 11 seconds | 5 seconds | 1004.45 MiB |
| 2020 | England/west-sussex | 100 | 375,837 | 871,029 | 337.76 MiB | 12 seconds | 5 seconds | 1.01 GiB |
| 2022 | England/west-sussex | 100 | 419,347 | 931,573 | 350.11 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2032 | England/west-sussex | 100 | 419,347 | 931,573 | 350.11 MiB | 12 seconds | 5 seconds | 1.03 GiB |
| 2039 | England/west-sussex | 100 | 442,292 | 958,567 | 356.77 MiB | 12 seconds | 5 seconds | 1.04 GiB |
| 2012 | England/west-yorkshire | 299 | 921,242 | 2,271,833 | 193.87 MiB | 46 seconds | 31 seconds | 3.05 GiB |
| 2020 | England/west-yorkshire | 299 | 963,460 | 2,339,930 | 190.47 MiB | 48 seconds | 33 seconds | 3.12 GiB |

| year | study_area | num_msaas | houses | households | people | file_size | runtime | commuting_time | memory_usage |
|------|--------------------------------|-----------|-----------|------------|------------|------------|------------|----------------|--------------|
| 2022 | England/west-yorkshire | 299 | 1,021,830 | 2,434,902 | 245.77 MiB | 48 seconds | 33 seconds | 3.13 GiB | |
| 2032 | England/west-yorkshire | 299 | 1,021,830 | 2,434,902 | 245.77 MiB | 49 seconds | 33 seconds | 3.13 GiB | |
| 2039 | England/west-yorkshire | 299 | 1,053,859 | 2,481,359 | 257.40 MiB | 49 seconds | 33 seconds | 3.32 GiB | |
| 2012 | England/wiltshire | 89 | 285,600 | 704,491 | 274.58 MiB | 9 seconds | 3 seconds | 921.08 MiB | |
| 2020 | England/wiltshire | 89 | 309,159 | 735,088 | 288.20 MiB | 9 seconds | 3 seconds | 947.43 MiB | |
| 2022 | England/wiltshire | 89 | 335,400 | 774,105 | 292.69 MiB | 10 seconds | 3 seconds | 949.16 MiB | |
| 2032 | England/wiltshire | 89 | 335,400 | 774,105 | 292.69 MiB | 10 seconds | 3 seconds | 949.16 MiB | |
| 2039 | England/wiltshire | 89 | 348,866 | 792,075 | 296.40 MiB | 10 seconds | 3 seconds | 955.08 MiB | |
| 2012 | England/worcestershire85 | 85 | 240,958 | 578,628 | 221.47 MiB | 8 seconds | 3 seconds | 770.62 MiB | |
| 2020 | England/worcestershire85 | 85 | 255,594 | 601,116 | 231.59 MiB | 8 seconds | 3 seconds | 790.42 MiB | |
| 2022 | England/worcestershire85 | 85 | 274,309 | 644,922 | 241.99 MiB | 8 seconds | 3 seconds | 849.84 MiB | |
| 2032 | England/worcestershire85 | 85 | 274,309 | 644,922 | 241.99 MiB | 8 seconds | 3 seconds | 849.84 MiB | |
| 2039 | England/worcestershire85 | 85 | 283,275 | 666,303 | 248.38 MiB | 8 seconds | 3 seconds | 861.38 MiB | |
| 2012 | special/northwest_transpennine | 829 | 2,653,096 | 6,416,492 | 2.45 GiB | 5 minutes | 4 minutes | 7.74 GiB | |
| 2020 | special/northwest_transpennine | 829 | 2,788,624 | 6,616,112 | 2.56 GiB | 5 minutes | 4 minutes | 7.95 GiB | |
| 2022 | special/northwest_transpennine | 829 | 2,960,285 | 6,908,374 | 2.62 GiB | 5 minutes | 4 minutes | 8.02 GiB | |

| year | study_area | num_msas | num_households | num_people | file_size | time | commuting_time | memory_usage |
|------|--------------------------------------|----------|----------------|------------|------------|------------|----------------|--------------|
| 2032 | special/northwest_transpennine | 829 | 2,960,285 | 6,908,374 | 2.62 GiB | 5 minutes | 5 minutes | 8.02 GiB |
| 2039 | special/northwest_transpennine | 829 | 3,058,114 | 7,059,122 | 2.66 GiB | 5 minutes | 5 minutes | 8.08 GiB |
| 2012 | Wales/bridgend-and-neath-port-talbot | 38 | 119,725 | 283,159 | 108.21 MiB | 5 seconds | 1 second | 382.24 MiB |
| 2020 | Wales/bridgend-and-neath-port-talbot | 38 | 123,909 | 289,896 | 111.10 MiB | 5 seconds | 1 second | 387.45 MiB |
| 2022 | Wales/bridgend-and-neath-port-talbot | 38 | 124,921 | 292,227 | 111.51 MiB | 4 seconds | 1 second | 387.72 MiB |
| 2032 | Wales/bridgend-and-neath-port-talbot | 38 | 128,601 | 301,529 | 113.58 MiB | 5 seconds | 1 second | 390.82 MiB |
| 2039 | Wales/bridgend-and-neath-port-talbot | 38 | 129,740 | 307,260 | 114.33 MiB | 5 seconds | 1 second | 391.29 MiB |
| 2012 | Wales/cardiff-and-vale-of-glamorgan | 63 | 199,208 | 484,182 | 187.17 MiB | 6 seconds | 2 seconds | 558.19 MiB |
| 2020 | Wales/cardiff-and-vale-of-glamorgan | 63 | 214,676 | 499,272 | 194.70 MiB | 7 seconds | 2 seconds | 572.89 MiB |
| 2022 | Wales/cardiff-and-vale-of-glamorgan | 63 | 218,981 | 502,763 | 196.11 MiB | 6 seconds | 2 seconds | 576.04 MiB |
| 2032 | Wales/cardiff-and-vale-of-glamorgan | 63 | 240,112 | 522,526 | 199.42 MiB | 7 seconds | 2 seconds | 577.84 MiB |
| 2039 | Wales/cardiff-and-vale-of-glamorgan | 63 | 254,162 | 531,549 | 201.82 MiB | 7 seconds | 2 seconds | 737.29 MiB |
| 2012 | Wales/central-valleys | 38 | 124,691 | 296,581 | 115.15 MiB | 5 seconds | 1 second | 396.20 MiB |
| 2020 | Wales/central-valleys | 38 | 130,072 | 301,907 | 117.77 MiB | 18 seconds | 1 second | 400.97 MiB |
| 2022 | Wales/central-valleys | 38 | 131,383 | 303,557 | 118.40 MiB | 8 seconds | 1 second | 424.47 MiB |
| 2032 | Wales/central-valleys | 38 | 136,404 | 310,032 | 118.04 MiB | 5 seconds | 1 second | 421.13 MiB |
| 2039 | Wales/central-valleys | 38 | 138,735 | 314,703 | 119.17 MiB | 5 seconds | 1 second | 423.02 MiB |
| 2012 | Wales/conwy-and-denbighshire | 30 | 92,732 | 211,205 | 80.50 MiB | 4 seconds | 1 second | 251.47 MiB |
| 2020 | Wales/conwy-and-denbighshire | 30 | 95,314 | 213,302 | 81.56 MiB | 4 seconds | 1 second | 253.63 MiB |

| year | study_area | num_msas | households | people | file_size | time | commuting | memory_usage |
|------|------------------------------|----------|------------|---------|------------|-----------|-----------|--------------|
| 2022 | Wales/conwy-and-denbighshire | 30 | 95,881 | 214,182 | 81.85 MiB | 4 seconds | 1 second | 254.21 MiB |
| 2032 | Wales/conwy-and-denbighshire | 30 | 97,683 | 218,122 | 81.11 MiB | 4 seconds | 1 second | 251.17 MiB |
| 2039 | Wales/conwy-and-denbighshire | 30 | 97,687 | 220,933 | 80.92 MiB | 4 seconds | 1 second | 249.76 MiB |
| 2012 | Wales/flintshire-and-wrexham | 38 | 122,180 | 288,696 | 113.32 MiB | 5 seconds | 1 second | 393.63 MiB |
| 2020 | Wales/flintshire-and-wrexham | 38 | 127,660 | 292,056 | 114.58 MiB | 5 seconds | 1 second | 395.27 MiB |
| 2022 | Wales/flintshire-and-wrexham | 38 | 129,007 | 292,644 | 115.03 MiB | 5 seconds | 1 second | 396.56 MiB |
| 2032 | Wales/flintshire-and-wrexham | 38 | 134,527 | 292,817 | 112.37 MiB | 5 seconds | 1 second | 410.92 MiB |
| 2039 | Wales/flintshire-and-wrexham | 38 | 136,425 | 293,540 | 112.22 MiB | 5 seconds | 1 second | 410.77 MiB |
| 2012 | Wales/gwent-valleys | 46 | 144,178 | 341,543 | 132.17 MiB | 5 seconds | 1 second | 451.03 MiB |
| 2020 | Wales/gwent-valleys | 46 | 148,386 | 344,566 | 132.83 MiB | 5 seconds | 1 second | 450.89 MiB |
| 2022 | Wales/gwent-valleys | 46 | 149,374 | 345,498 | 132.72 MiB | 5 seconds | 1 second | 450.23 MiB |
| 2032 | Wales/gwent-valleys | 46 | 151,842 | 347,976 | 130.50 MiB | 5 seconds | 1 second | 442.86 MiB |
| 2039 | Wales/gwent-valleys | 46 | 151,729 | 350,397 | 130.58 MiB | 5 seconds | 1 second | 443.03 MiB |
| 2012 | Wales/gwynedd | 17 | 52,926 | 122,595 | 48.28 MiB | 3 seconds | 1 second | 141.47 MiB |
| 2020 | Wales/gwynedd | 17 | 55,064 | 124,569 | 49.28 MiB | 3 seconds | 1 second | 143.70 MiB |
| 2022 | Wales/gwynedd | 17 | 55,683 | 125,030 | 49.20 MiB | 3 seconds | 1 second | 143.45 MiB |
| 2032 | Wales/gwynedd | 17 | 58,372 | 128,844 | 49.81 MiB | 3 seconds | 1 second | 143.80 MiB |
| 2039 | Wales/gwynedd | 17 | 59,746 | 130,948 | 50.64 MiB | 3 seconds | 1 second | 145.62 MiB |
| 2012 | Wales/isle-of-anglesey | 9 | 30,797 | 69,919 | 27.65 MiB | 3 seconds | 1 second | 96.81 MiB |
| 2020 | Wales/isle-of-anglesey | 9 | 31,366 | 69,845 | 27.85 MiB | 3 seconds | 1 second | 97.39 MiB |

| year | study_area | num_msoas | houses | households | people | file_size | runtime | commuting_time | memory_usage |
|------|---------------------------------|-----------|---------|------------|---------------|----------------|----------|----------------|--------------|
| 2022 | Wales/isle-of-anglesey | 9 | 31,488 | 69,864 | 27.91 MiB | 3 sec- onds | 1 second | 97.71 MiB | |
| 2032 | Wales/isle-of-anglesey | 9 | 31,601 | 69,502 | 27.09 MiB | 3 sec- onds | 1 second | 95.51 MiB | |
| 2039 | Wales/isle-of-anglesey | 9 | 31,337 | 69,423 | 26.91 MiB | 3 sec- onds | 1 second | 95.37 MiB | |
| 2012 | Wales/monmouthshire-and-newport | 31 | 100,402 | 240,491 | 94.44 MiB | 4 sec- onds | 1 second | 280.40 MiB | |
| 2020 | Wales/monmouthshire-and-newport | 31 | 104,394 | 250,185 | 98.11 MiB | 4 sec- onds | 1 second | 286.97 MiB | |
| 2022 | Wales/monmouthshire-and-newport | 31 | 105,481 | 253,282 | 99.27 MiB | 4 sec- onds | 1 second | 289.03 MiB | |
| 2032 | Wales/monmouthshire-and-newport | 31 | 109,752 | 265,785 | 102.21 MiB | 4 sec- onds | 1 second | 371.39 MiB | |
| 2039 | Wales/monmouthshire-and-newport | 31 | 111,246 | 273,319 | 103.90 MiB | 4 sec- onds | 1 second | 373.82 MiB | |
| 2012 | Wales/powys | 19 | 59,028 | 132,725 | 51.21 MiB | 4 sec- onds | 1 second | 185.06 MiB | |
| 2020 | Wales/powys | 19 | 59,972 | 132,328 | 50.60 MiB | 4 sec- onds | 1 second | 183.38 MiB | |
| 2022 | Wales/powys | 19 | 60,190 | 132,467 | 50.46 MiB | 4 sec- onds | 1 second | 182.88 MiB | |
| 2032 | Wales/powys | 19 | 59,586 | 133,010 | 49.63 MiB | 4 sec- onds | 1 second | 180.64 MiB | |
| 2039 | Wales/powys | 19 | 57,969 | 133,514 | 49.35 MiB | 4 sec- onds | 1 second | 179.80 MiB | |
| 2012 | Wales/south-west-wales | 50 | 165,004 | 383,260 | 145.79 MiB | 5 sec- onds | 1 second | 474.32 MiB | |
| 2020 | Wales/south-west-wales | 50 | 170,327 | 385,937 | 146.53 MiB | 5 sec- onds | 1 second | 474.47 MiB | |
| 2022 | Wales/south-west-wales | 50 | 171,623 | 386,901 | 147.00 MiB | 5 sec- onds | 1 second | 476.10 MiB | |
| 2032 | Wales/south-west-wales | 50 | 175,897 | 392,107 | 145.19 MiB | 6 sec- onds | 1 second | 469.31 MiB | |
| 2039 | Wales/south-west-wales | 50 | 176,482 | 394,303 | 144.53 MiB | 6 sec- onds | 1 second | 467.47 MiB | |
| 2012 | Wales/swansea | 31 | 104,423 | 242,128 | 93.10 MiB | 4 sec- onds | 1 second | 276.14 MiB | |
| 2020 | Wales/swansea | 31 | 110,304 | 247,820 | 95.72 MiB | 4 sec- onds | 1 second | 281.38 MiB | |

| year | study_area | num_msas | num_households | num_people | pb_file_size | runtime | commuting_runtime | memory_usage |
|------|---------------|----------|----------------|------------|--------------|-----------|-------------------|--------------|
| 2022 | Wales/swansea | 31 | 111,940 | 249,098 | 96.10 MiB | 4 seconds | 1 second | 282.16 MiB |
| 2032 | Wales/swansea | 31 | 119,141 | 257,653 | 98.28 MiB | 4 seconds | 1 second | 285.53 MiB |
| 2039 | Wales/swansea | 31 | 123,450 | 262,306 | 99.93 MiB | 4 seconds | 1 second | 366.61 MiB |

Notes:

- `pb_file_size` refers to the size of the uncompressed protobuf file in `data/output/`
- The total `runtime` is usually dominated by matching workers to businesses, so `commuting_runtime` gives a breakdown
- Measuring memory usage of Linux processes isn't straightforward, so `memory_usage` should just be a guide
- These measurements were all taken on one developer's laptop, and they don't represent multiple runs. This table just aims to give a general sense of how long running takes.
 - That machine has 24 cores, which matters for the parallelized commuting calculation.
- The time *usually* doesn't include downloading or decompressing raw data. For some areas, it might!
- `scripts/collect_stats.py` produces the table above